

Reward Shaping, Advanced Exploration, and Entropy Regularization

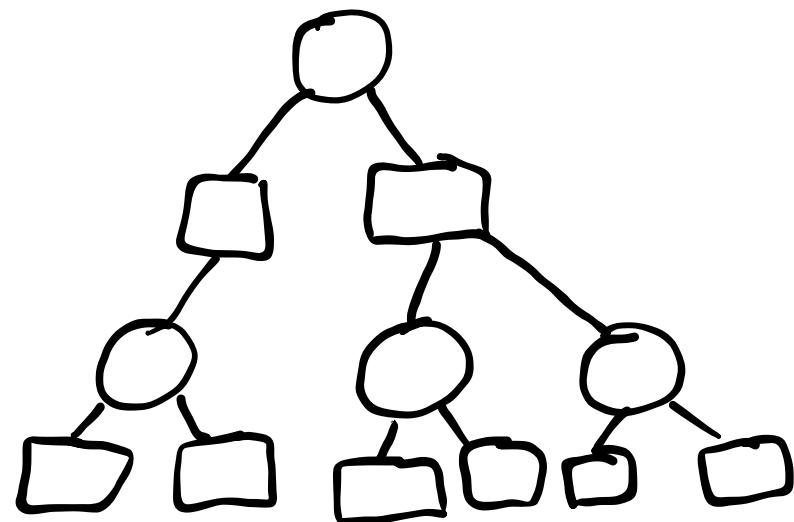
Alpha Zero: Actor Critic with MCTS

Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS

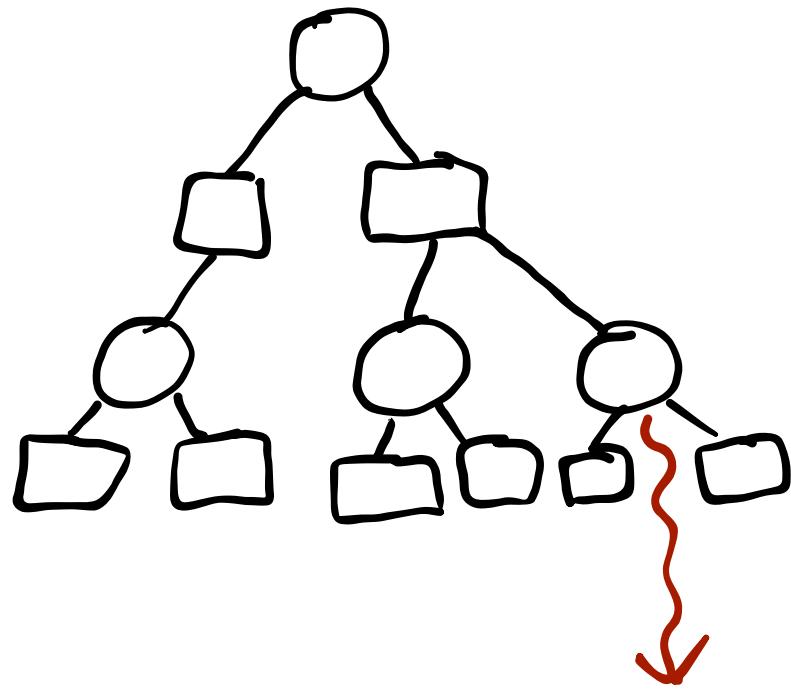
Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS



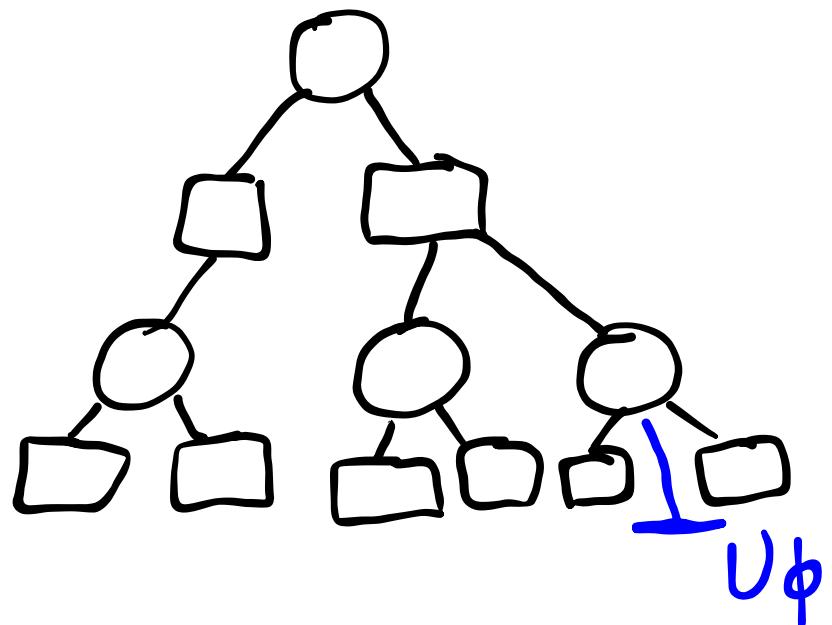
Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS



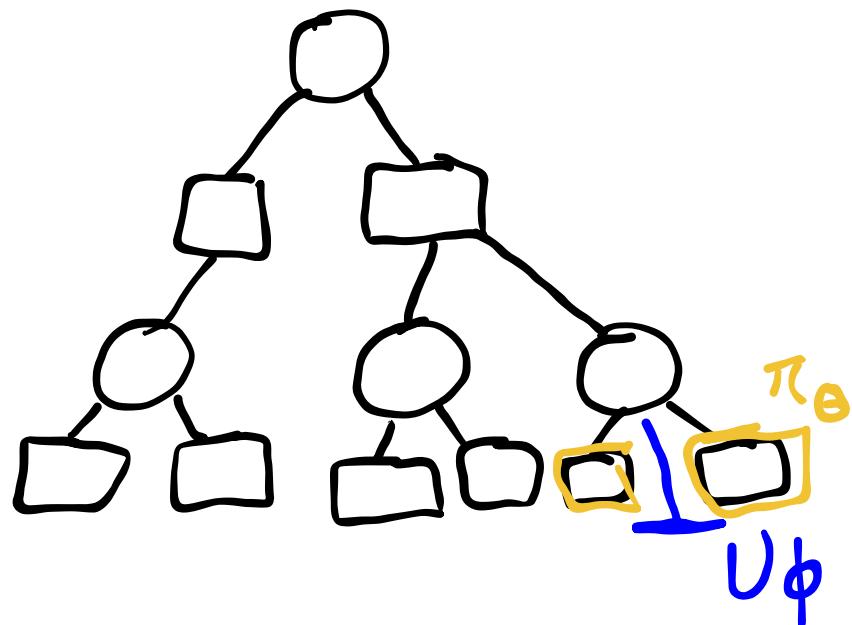
Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS



Alpha Zero: Actor Critic with MCTS

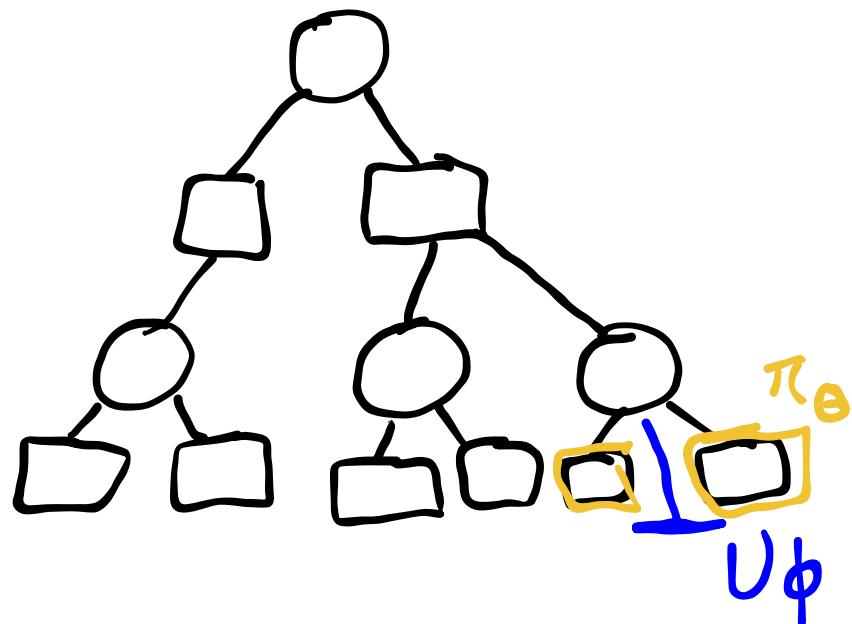
1. Use π_θ and U_ϕ in MCTS



$$a = \arg \max_a Q(s, a) + c\pi_\theta(a \mid s) \frac{\sqrt{N(s)}}{1 + N(s, a)}$$

Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS
2. Learn π_θ and U_ϕ from tree

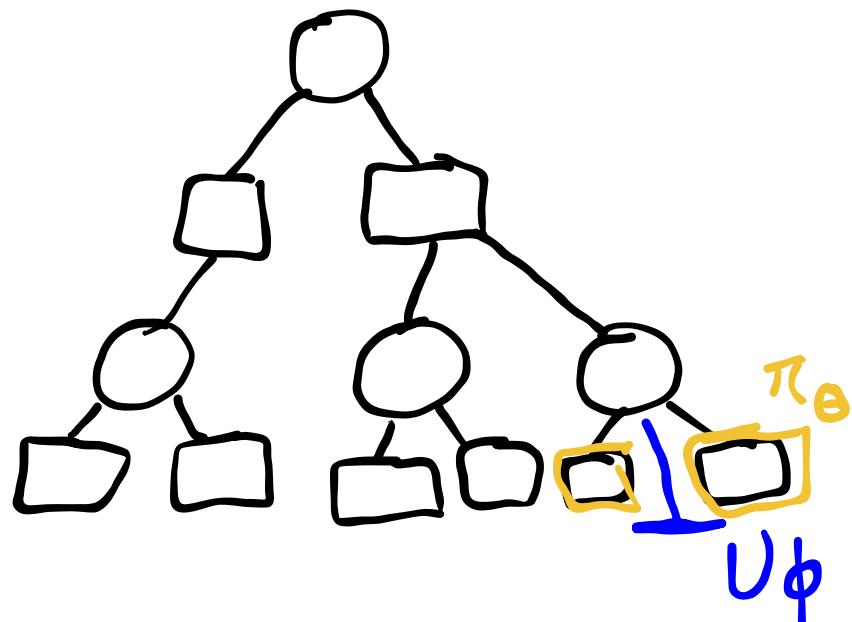


$$a = \arg \max_a Q(s, a) + c\pi_\theta(a \mid s) \frac{\sqrt{N(s)}}{1 + N(s, a)}$$

Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS
2. Learn π_θ and U_ϕ from tree

$$\ell(\boldsymbol{\theta}) = -\mathbb{E}_s \left[\sum_a \pi_{\text{MCTS}}(a | s) \log \pi_\theta(a | s) \right]$$
$$\pi_{\text{MCTS}}(a | s) \propto N(s, a)^\eta$$

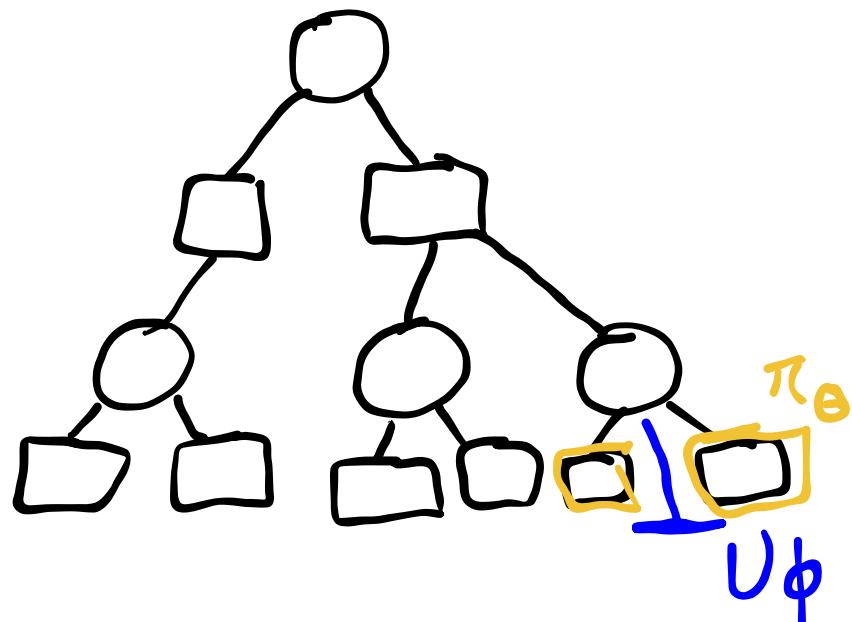


$$a = \arg \max_a Q(s, a) + c \pi_\theta(a | s) \frac{\sqrt{N(s)}}{1 + N(s, a)}$$

Alpha Zero: Actor Critic with MCTS

1. Use π_θ and U_ϕ in MCTS
2. Learn π_θ and U_ϕ from tree

$$\ell(\boldsymbol{\theta}) = -\mathbb{E}_s \left[\sum_a \pi_{\text{MCTS}}(a | s) \log \pi_\theta(a | s) \right]$$
$$\pi_{\text{MCTS}}(a | s) \propto N(s, a)^\eta$$



$$\ell(\boldsymbol{\Phi}) = \frac{1}{2} \mathbb{E}_s \left[(U_\Phi(s) - U_{\text{MCTS}}(s))^2 \right]$$

$$U_{\text{MCTS}}(s) = \max_a Q(s, a)$$

$$a = \arg \max_a Q(s, a) + c \pi_\theta(a | s) \frac{\sqrt{N(s)}}{1 + N(s, a)}$$

<https://www.youtube.com/embed/tI0IHko8ySg?enablejsapi=1>

<https://www.youtube.com/watch?v=tI0IHko8ySg>

Reward Shaping

"As a general rule, it is better to design performance measures according to what one actually wants in the environment, rather than according to how one thinks the agent should behave." - Stuart Russell

Reward Shaping

"As a general rule, it is better to design performance measures according to what one actually wants in the environment, rather than according to how one thinks the agent should behave." - Stuart Russell

Reward

| | | | | | | | | | |
|------|------|------|------|------|------|------|------|------|------|
| -0.2 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.2 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 3 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | -5 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | -10 | 0 | 0 | 0 | 10 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.2 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.2 |

Reward Shaping

"As a general rule, it is better to design performance measures according to what one actually wants in the environment, rather than according to how one thinks the agent should behave." - Stuart Russell

Reward

| | | | | | | | | | |
|------|------|------|------|------|------|------|------|------|------|
| -0.2 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.2 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 3 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | -0.1 |
| -0.2 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.1 | -0.2 |

Value

| | | | | | | | | | |
|------|------|------|-------|------|------|------|------|------|------|
| 0.41 | 0.74 | 0.96 | 1.18 | 1.43 | 1.71 | 1.98 | 2.11 | 2.39 | 2.09 |
| 0.74 | 1.04 | 1.27 | 1.52 | 1.81 | 2.15 | 2.47 | 2.58 | 3.02 | 2.69 |
| 0.86 | 1.18 | 1.45 | 1.76 | 2.15 | 2.55 | 2.97 | 3 | 3.69 | 3.32 |
| 0.84 | 1.11 | 1.31 | 1.55 | 2.45 | 3.01 | 3.56 | 4.1 | 4.53 | 4.04 |
| 0.91 | 1.2 | 1.09 | -3 | 2.48 | 3.53 | 4.21 | 4.93 | 5.5 | 4.88 |
| 1.1 | 1.46 | 1.79 | 2.24 | 3.42 | 4.2 | 4.97 | 5.85 | 6.68 | 5.84 |
| 1.06 | 1.41 | 1.7 | 2.14 | 3.89 | 4.9 | 5.85 | 6.92 | 8.15 | 6.94 |
| 0.92 | 1.18 | 0.7 | -7.39 | 3.43 | 5.39 | 6.67 | 8.15 | 10 | 8.19 |
| 1.09 | 1.45 | 1.75 | 2.18 | 3.89 | 4.88 | 5.84 | 6.92 | 8.15 | 6.94 |
| 1.07 | 1.56 | 2.05 | 2.65 | 3.38 | 4.11 | 4.92 | 5.83 | 6.68 | 5.82 |

Reward Shaping

Potential-Based Reward Shaping

$$\gamma \underline{\phi(s')} - \underline{\phi(s)}$$

- $\underline{R(s, a, s')} + = \underline{F(s)} - \gamma \underline{F(s')}$
- any other transformation may yield sub optimal policies unless further assumptions are made about the underlying MDP

Is Exploration Important? Montezuma's Revenge

Is Exploration Important?

Theory

| | Algorithm | Regret | Time | Space |
|-------------|--|--|-------------------|------------------------|
| Model-based | UCRL2 [10] ¹ | at least $\tilde{\mathcal{O}}(\sqrt{H^4 S^2 A T})$ | $\Omega(T S^2 A)$ | $\mathcal{O}(S^2 A H)$ |
| | Agrawal and Jia [1] ¹ | at least $\tilde{\mathcal{O}}(\sqrt{H^3 S^2 A T})$ | | |
| | UCBVI [5] ² | $\tilde{\mathcal{O}}(\sqrt{H^2 S A T})$ | | |
| | vUCQ [12] ² | $\tilde{\mathcal{O}}(\sqrt{H^2 S A T})$ | | |
| Model-free | Q-learning (ε -greedy) [14] (if 0 initialized) | $\Omega(\min\{T, A^{H/2}\})$ | $\mathcal{O}(T)$ | $\mathcal{O}(S A H)$ |
| | Delayed Q-learning [25] ³ | $\tilde{\mathcal{O}}_{S,A,H}(T^{4/5})$ | | |
| | Q-learning (UCB-H) | $\tilde{\mathcal{O}}(\sqrt{H^4 S A T})$ | | |
| | Q-learning (UCB-B) | $\tilde{\mathcal{O}}(\sqrt{H^3 S A T})$ | | |
| | lower bound | $\Omega(\sqrt{H^2 S A T})$ | - | - |

Table 1: Regret comparisons for RL algorithms on episodic MDP. $T = KH$ is totally number of steps, H is the number of steps per episode, S is the number of states, and A is the number of actions. For clarity, this table is presented for $T \geq \text{poly}(S, A, H)$, omitting low order terms.

Exploration Bonus

Exploration Bonus

- In General, $R^+(s, a) = R(s, a) + B(s, a)$
- UCB: $B(s, a) = c\sqrt{\frac{\log N(s)}{N(s,a)}}$

Exploration Bonus

Example 1: Learn Pseudocount

Exploration Bonus

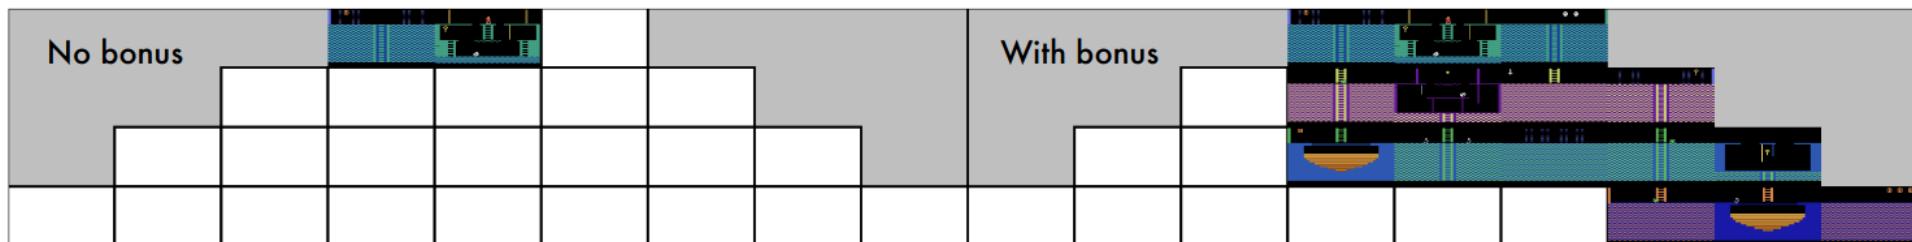
Example 1: Learn Pseudocount

$B(s, a) \approx \frac{1}{\sqrt{\hat{N}(s)}}$ where $\hat{N}(s)$ is a learned function approximation

Exploration Bonus

Example 1: Learn Pseudocount

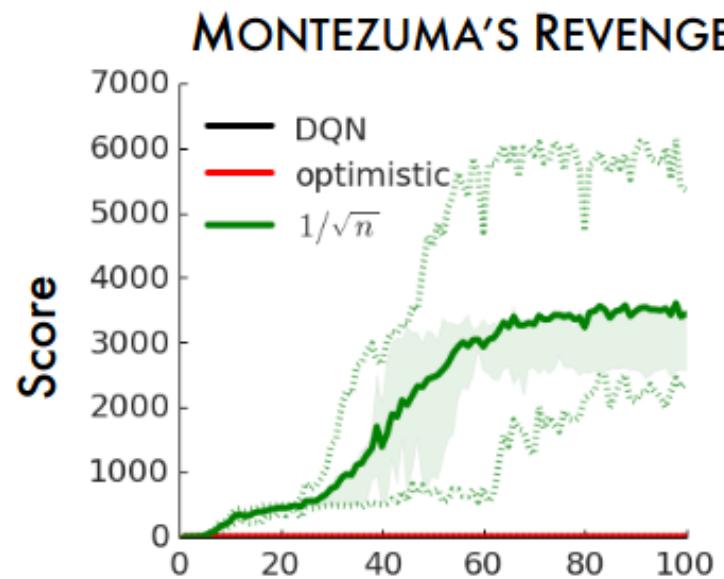
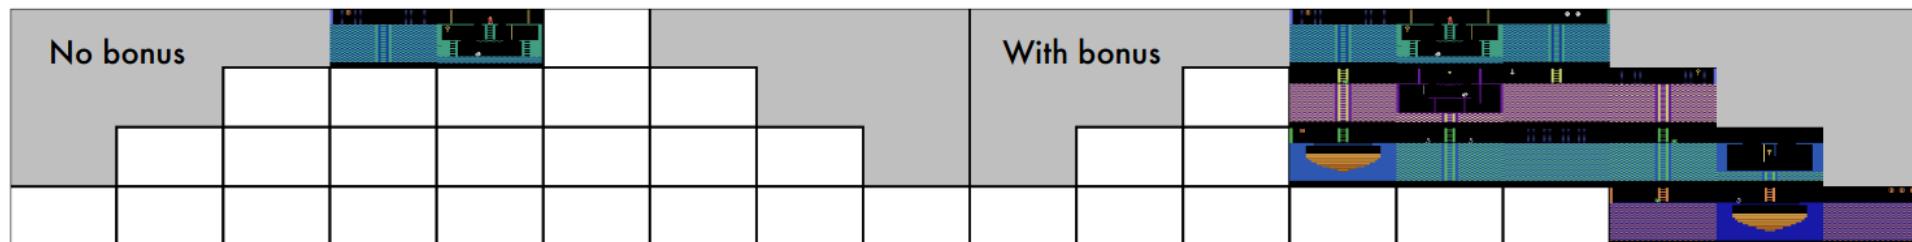
$$B(s, a) \approx \frac{1}{\sqrt{\hat{N}(s)}} \text{ where } \hat{N}(s) \text{ is a learned function approximation}$$



Exploration Bonus

Example 1: Learn Pseudocount

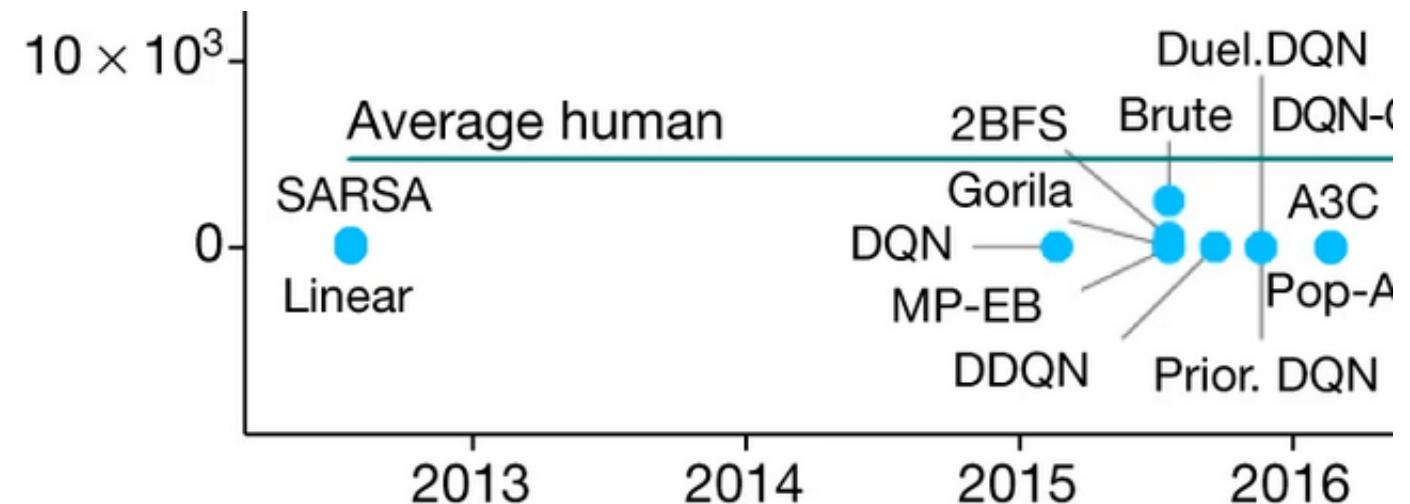
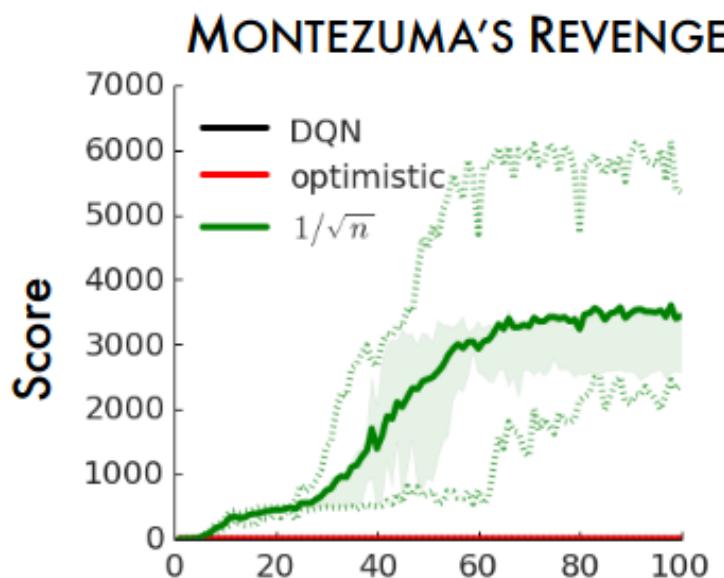
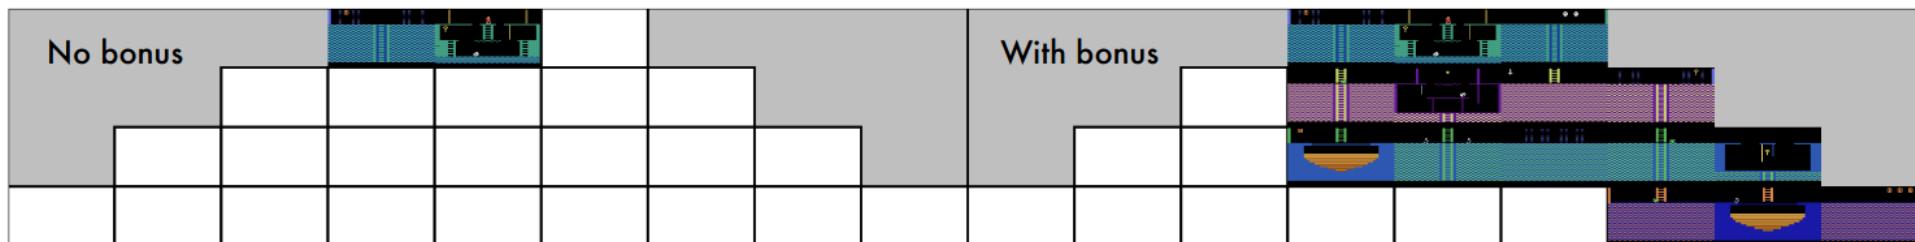
$B(s, a) \approx \frac{1}{\sqrt{\hat{N}(s)}}$ where $\hat{N}(s)$ is a learned function approximation



Exploration Bonus

Example 1: Learn Pseudocount

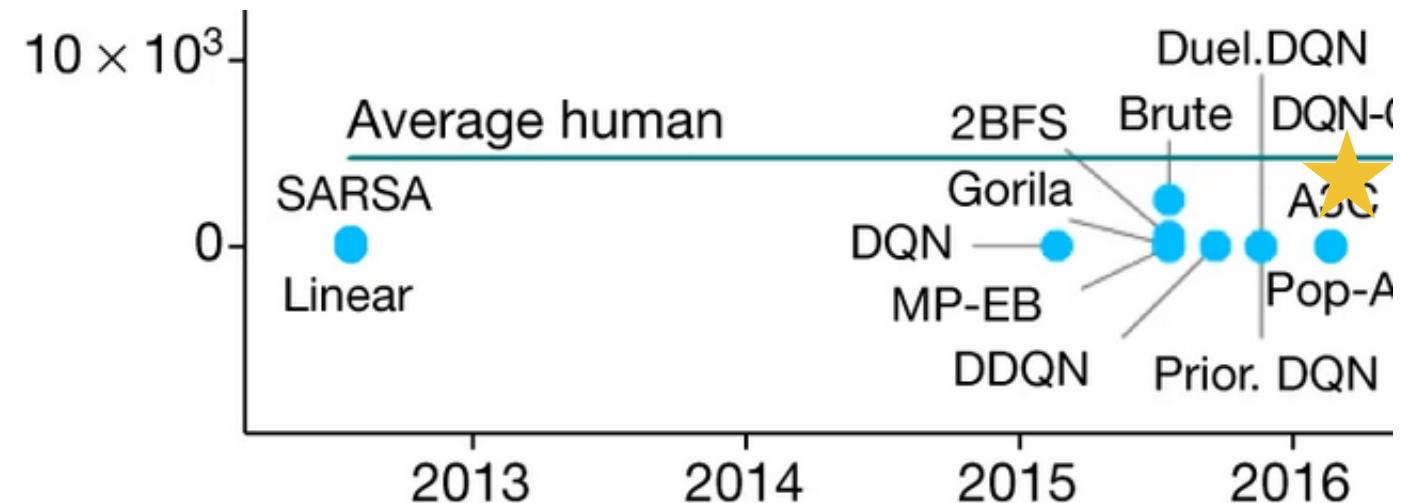
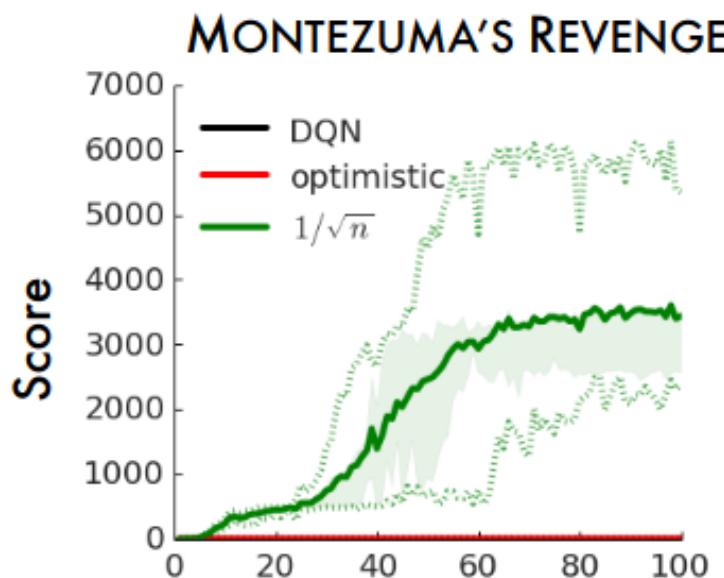
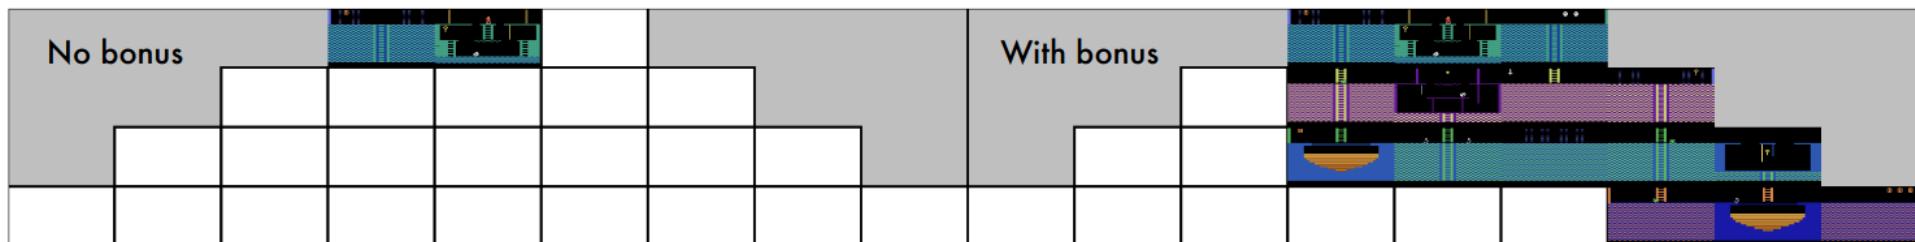
$B(s, a) \approx \frac{1}{\sqrt{\hat{N}(s)}}$ where $\hat{N}(s)$ is a learned function approximation



Exploration Bonus

Example 1: Learn Pseudocount

$B(s, a) \approx \frac{1}{\sqrt{\hat{N}(s)}}$ where $\hat{N}(s)$ is a learned function approximation



Exploration Bonus

^

-

,

Exploration Bonus

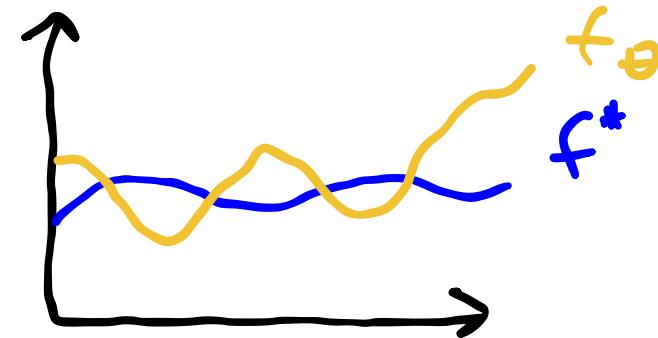
Example 2: Learn a function of the state and action



Exploration Bonus

Example 2: Learn a function of the state and action

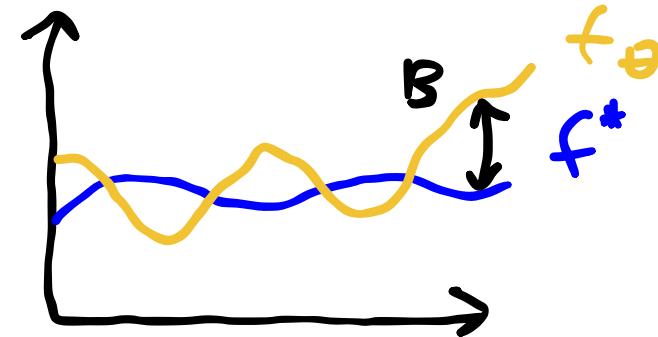
$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$



Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

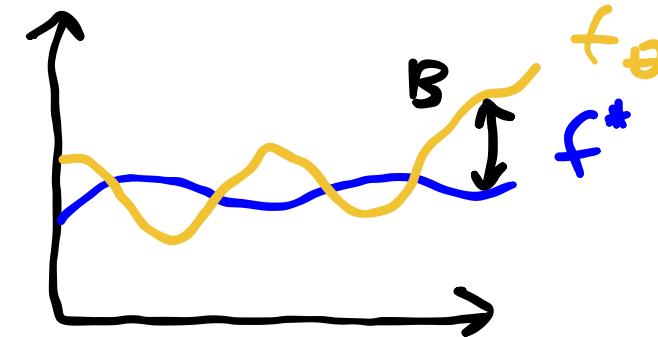


Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

What should f^* be?

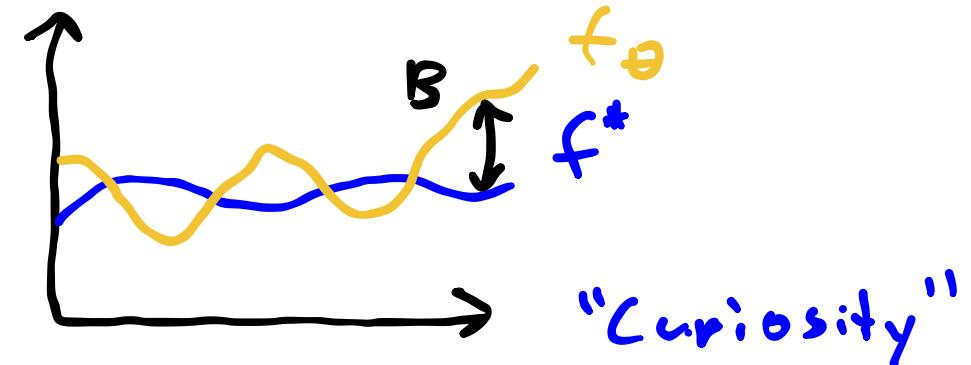


Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

What should f^* be?



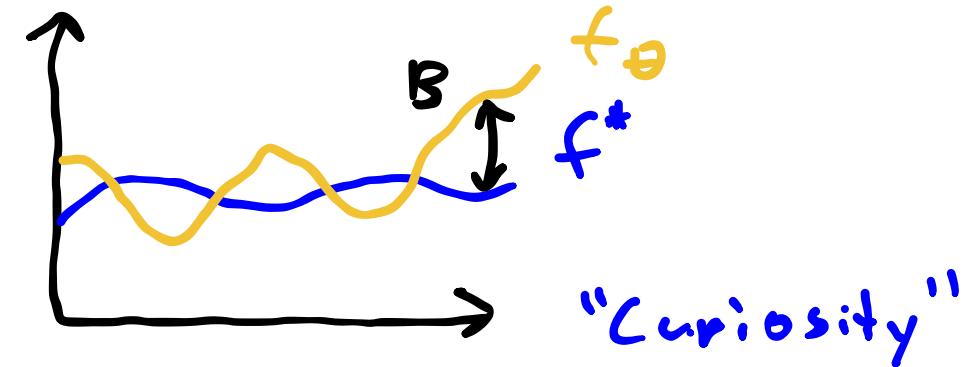
Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

What should f^* be?

- $f^*(s, a) = s'$ (Next state prediction)



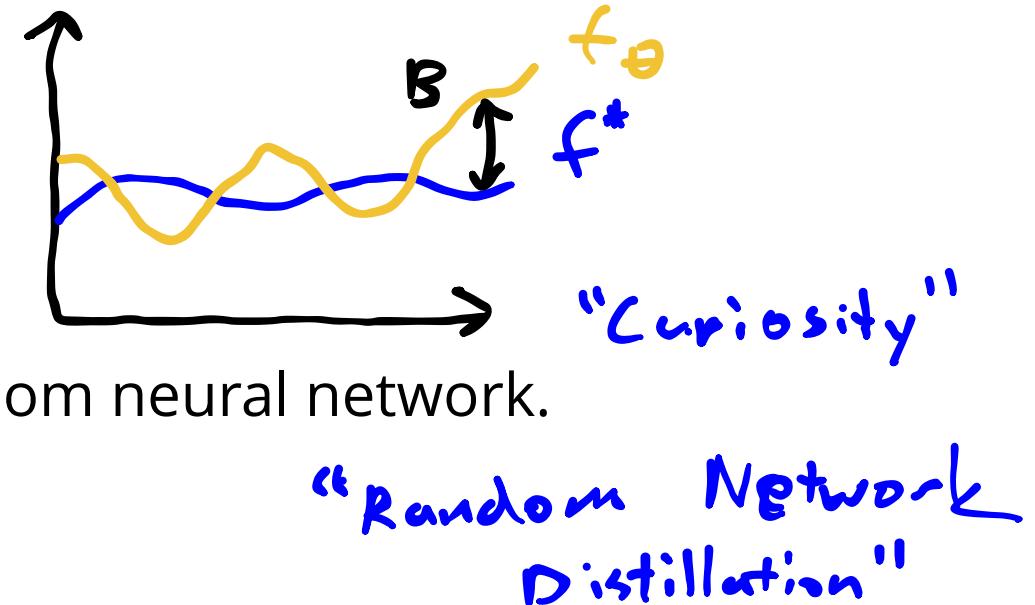
Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

What should f^* be?

- $f^*(s, a) = s'$ (Next state prediction)
- $f^*(s, a) = f_\phi(s, a)$ where f_ϕ is a random neural network.



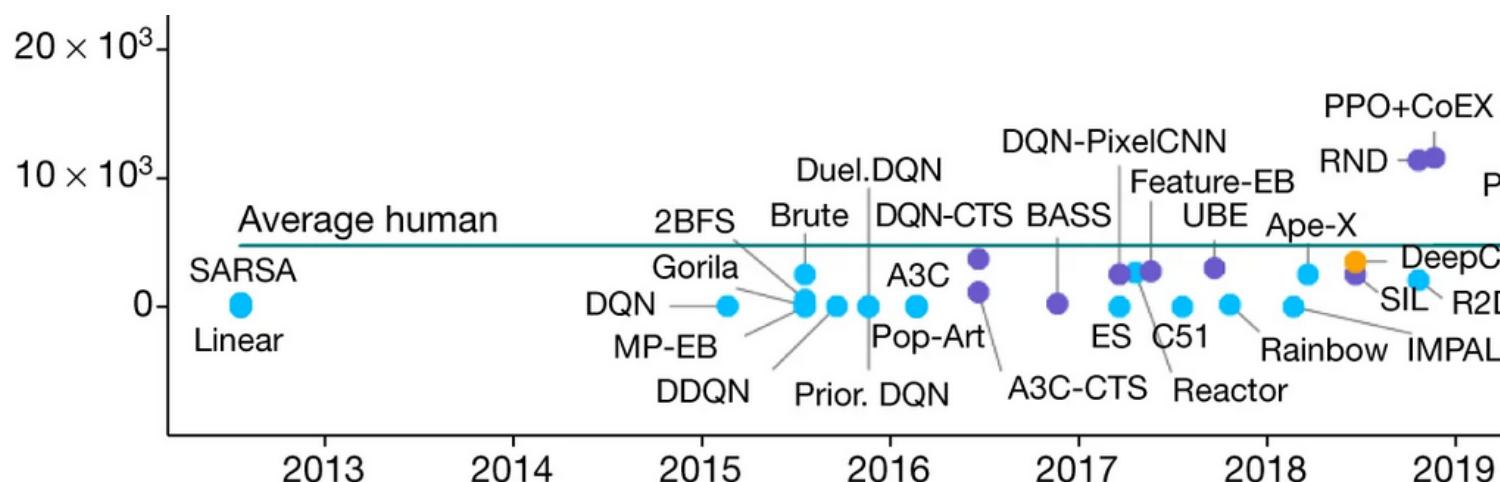
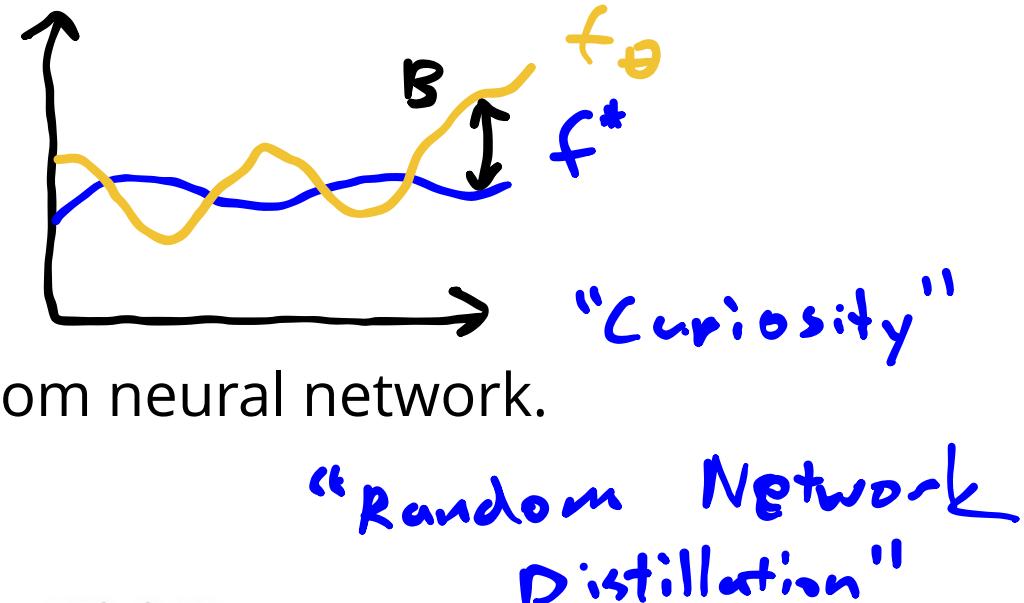
Exploration Bonus

Example 2: Learn a function of the state and action

$$B(s, a) = \|\hat{f}_\theta(s, a) - f^*(s, a)\|^2$$

What should f^* be?

- $f^*(s, a) = s'$ (Next state prediction)
- $f^*(s, a) = f_\phi(s, a)$ where f_ϕ is a random neural network.



Exploration Bonus

Exploration Bonus

Example 3: Thompson Sampling

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q
2. Sample Q

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q
2. Sample Q
3. Act according to Q

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q  Hard
2. Sample Q
3. Act according to Q

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q  Hard
 2. Sample Q
 3. Act according to Q
-
- Bootstrapping with multiple Q networks

Exploration Bonus

Example 3: Thompson Sampling

1. Maintain a distribution over Q  Hard
2. Sample Q
3. Act according to Q

- Bootstrapping with multiple Q networks
- Dropout

Exploration Bonus

Exploration Bonus

Example 4: Go-Explore

Exploration Bonus

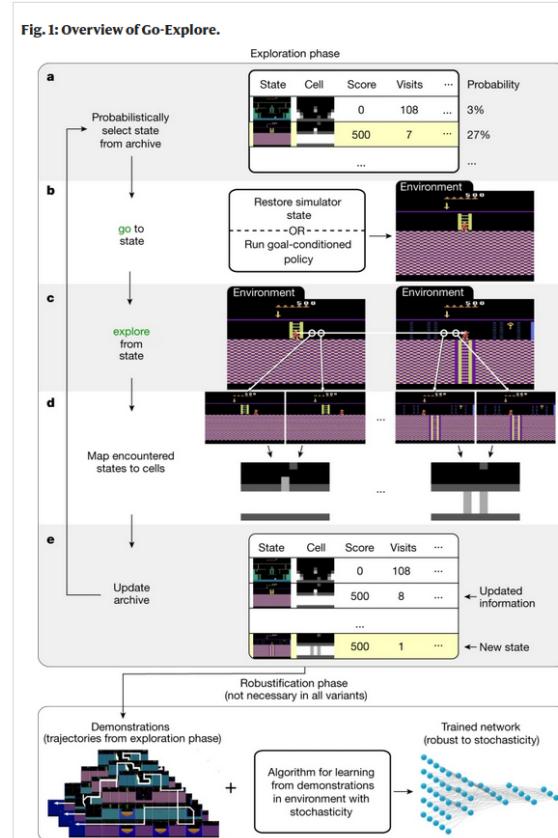
Example 4: Go-Explore

"First return, then explore"

Exploration Bonus

Example 4: Go-Explore

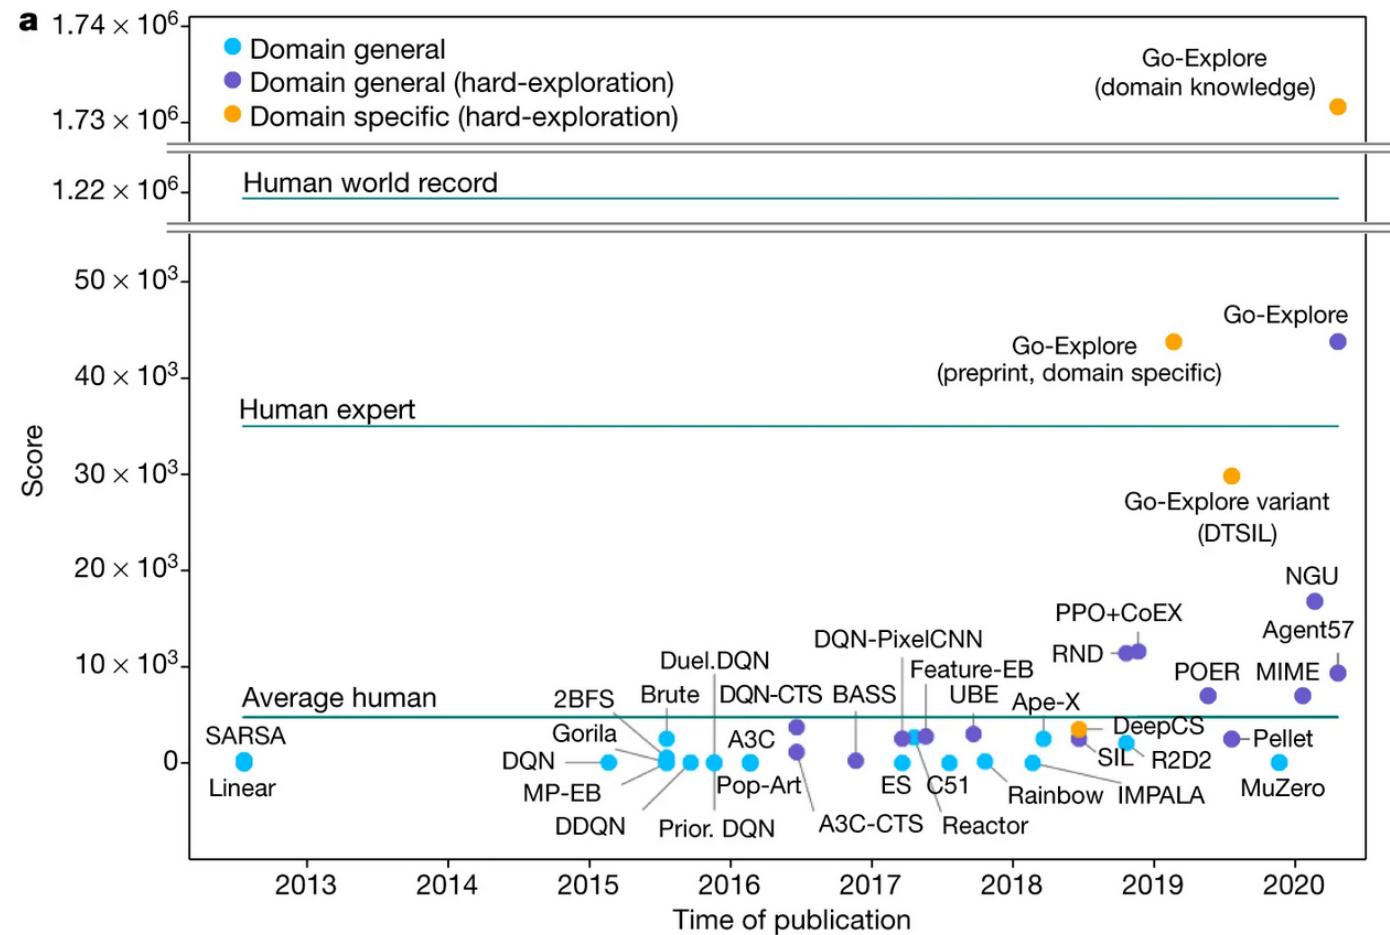
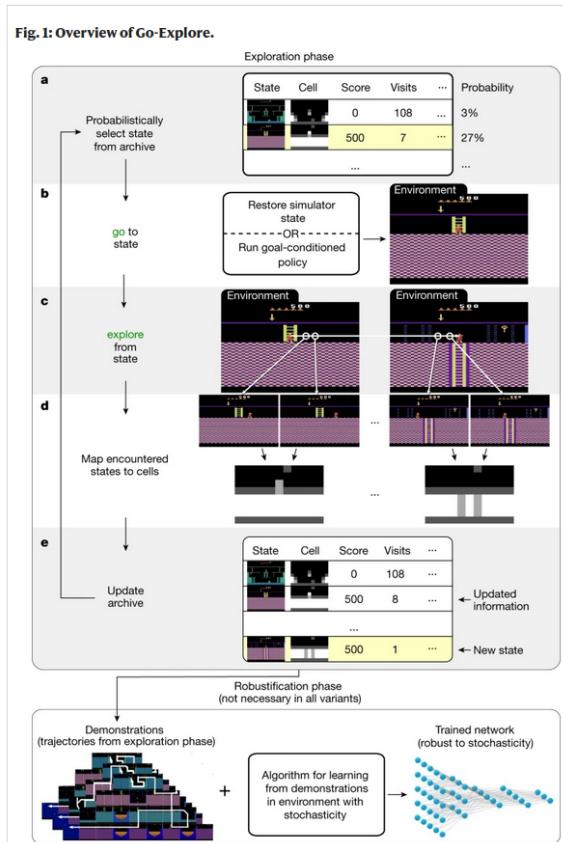
"First return, then explore"



Exploration Bonus

Example 4: Go-Explore

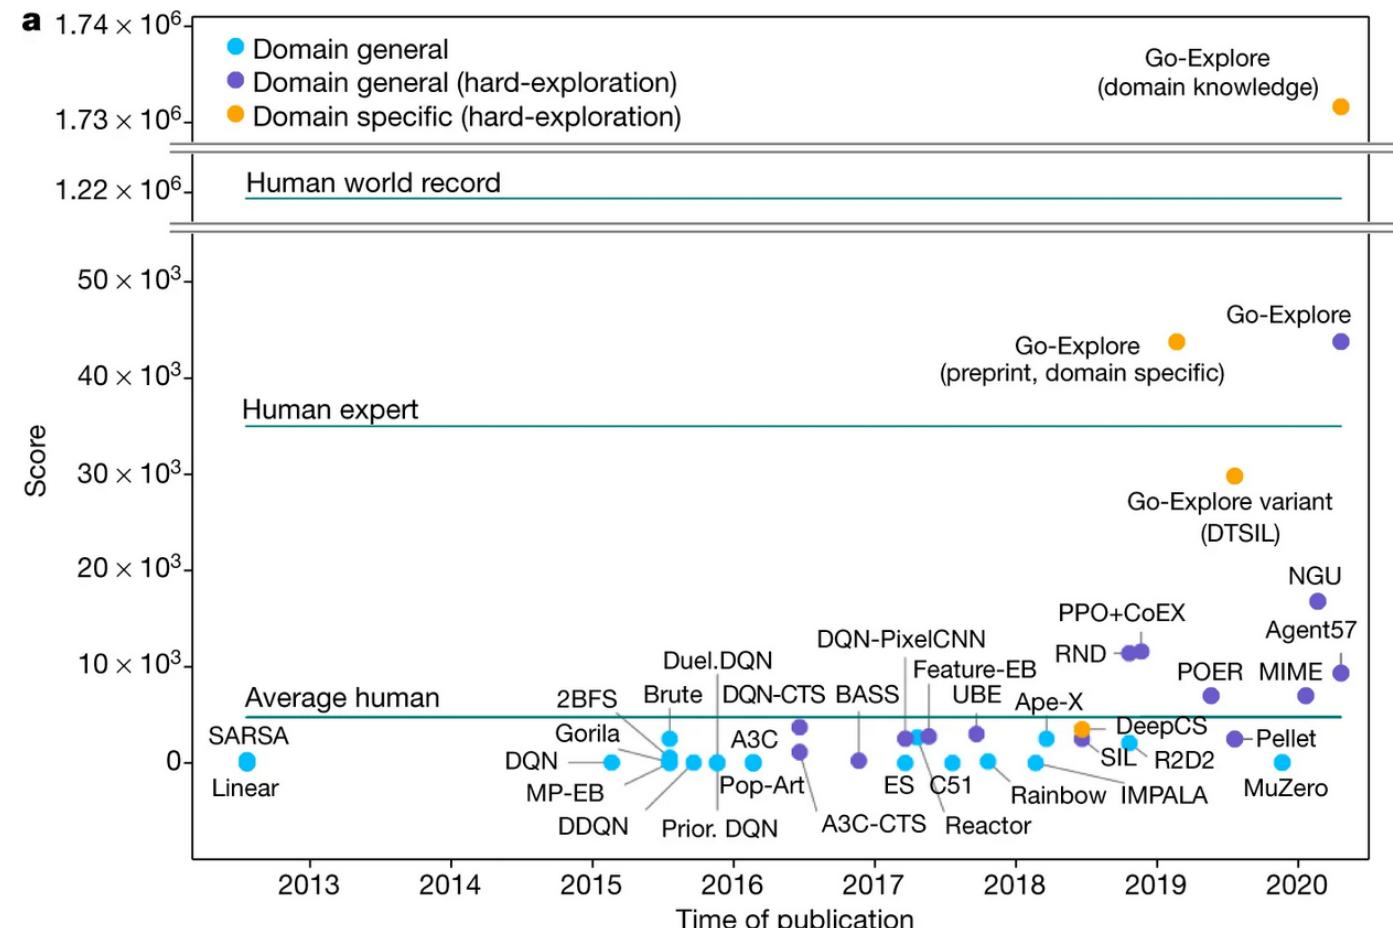
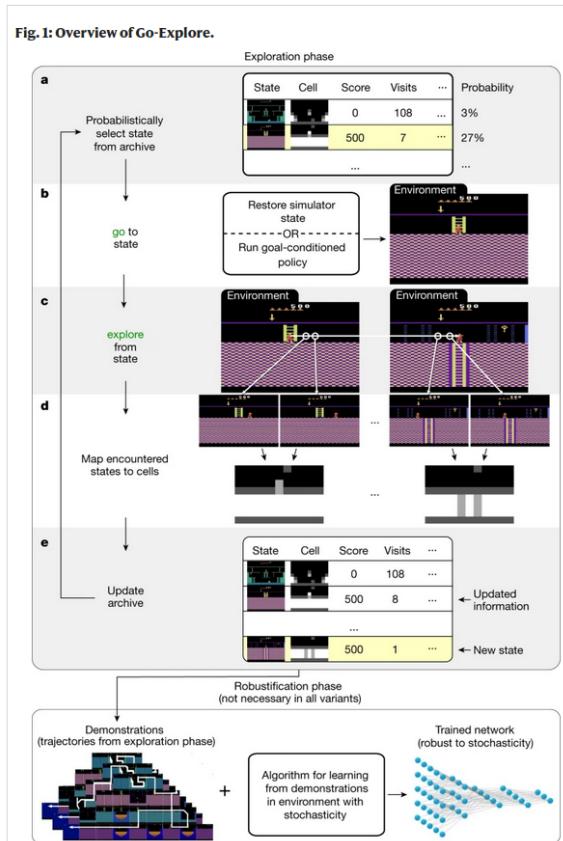
"First return, then explore"



Exploration Bonus

Example 4: Go-Explore

"First return, then explore"

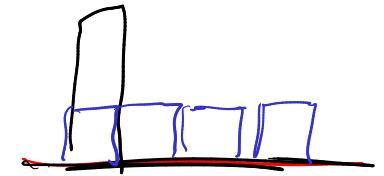


(Uber AI Labs)

Soft Actor Critic: Entropy Regularization

Soft Actor Critic: Entropy Regularization

$$U(\pi) = E \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha \mathcal{H}(\pi(\cdot | s_t))) \right]$$



Soft Actor Critic: Entropy Regularization

$$U(\pi) = E \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha \mathcal{H}(\pi(\cdot | s_t))) \right]$$

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(\mathbf{s}_t, \mathbf{a}_t) - \underbrace{\log \pi(\mathbf{a}_t | \mathbf{s}_t)}_{-}]$$

Soft Actor Critic: Entropy Regularization

$$U(\pi) = E \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha \mathcal{H}(\pi(\cdot | s_t))) \right]$$

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(\mathbf{s}_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t | \mathbf{s}_t)]$$

$$\mathcal{T}^\pi Q(\mathbf{s}_t, \mathbf{a}_t) \triangleq r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V(\mathbf{s}_{t+1})]$$

Soft Actor Critic: Entropy Regularization

$$U(\pi) = E \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha \mathcal{H}(\pi(\cdot | s_t))) \right]$$

$$V(s_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(s_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t | s_t)]$$

$$\xrightarrow{\text{underline}} \mathcal{T}^\pi Q(s_t, \mathbf{a}_t) \triangleq r(s_t, \mathbf{a}_t) + \gamma \mathbb{E}_{s_{t+1} \sim p} [V(s_{t+1})]$$

$$\xrightarrow{\text{underbrace}} \pi_{\text{new}} = \arg \min_{\substack{\pi' \in \Pi}} D_{\text{KL}} \left(\pi'(\cdot | s_t) \parallel \frac{\exp(Q^{\pi_{\text{old}}}(s_t, \cdot))}{Z^{\pi_{\text{old}}}(s_t)} \right)$$

Soft Actor Critic

Soft Actor Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau\psi + (1 - \tau)\bar{\psi}$$

end for

end for

Soft Actor Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

$$J_V(\psi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\frac{1}{2} \left(V_\psi(\mathbf{s}_t) - \mathbb{E}_{\mathbf{a}_t \sim \pi_\phi} [Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \log \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)] \right)^2 \right]$$

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau \psi + (1 - \tau) \bar{\psi}$$

end for

end for

Soft Actor Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

$$J_V(\psi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\frac{1}{2} \left(V_\psi(\mathbf{s}_t) - \mathbb{E}_{\mathbf{a}_t \sim \pi_\phi} [Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \log \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)] \right)^2 \right]$$

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau \psi + (1 - \tau) \bar{\psi}$$

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \underbrace{\hat{Q}(\mathbf{s}_t, \mathbf{a}_t)}_{\text{underlined}} \right)^2 \right]$$

end for

end for

Soft Actor Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

$$J_V(\psi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\frac{1}{2} \left(V_\psi(\mathbf{s}_t) - \mathbb{E}_{\mathbf{a}_t \sim \pi_\phi} [Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \log \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)] \right)^2 \right]$$

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau \psi + (1 - \tau) \bar{\psi}$$

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \hat{Q}(\mathbf{s}_t, \mathbf{a}_t) \right)^2 \right]$$

$$\hat{Q}(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V_{\bar{\psi}}(\mathbf{s}_{t+1})]$$

end for

end for

Soft Actor Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

$$J_V(\psi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\frac{1}{2} \left(V_\psi(\mathbf{s}_t) - \mathbb{E}_{\mathbf{a}_t \sim \pi_\phi} [Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \log \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)] \right)^2 \right]$$

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau \psi + (1 - \tau) \bar{\psi}$$

end for

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \hat{Q}(\mathbf{s}_t, \mathbf{a}_t) \right)^2 \right]$$

$$\hat{Q}(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V_{\bar{\psi}}(\mathbf{s}_{t+1})]$$

$$J_\pi(\phi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\text{D}_{\text{KL}} \left(\pi_\phi(\cdot | \mathbf{s}_t) \parallel \frac{\exp(Q_\theta(\mathbf{s}_t, \cdot))}{Z_\theta(\mathbf{s}_t)} \right) \right]$$

end for

Soft Actor Critic

Advantages:

Soft Actor Critic

Advantages:

- Stable

Soft Actor Critic

Advantages:

- Stable
- Learns many near-optimal policies

Soft Actor Critic

Advantages:

- Stable
- Learns many near-optimal policies
- Exploration

Soft Actor Critic

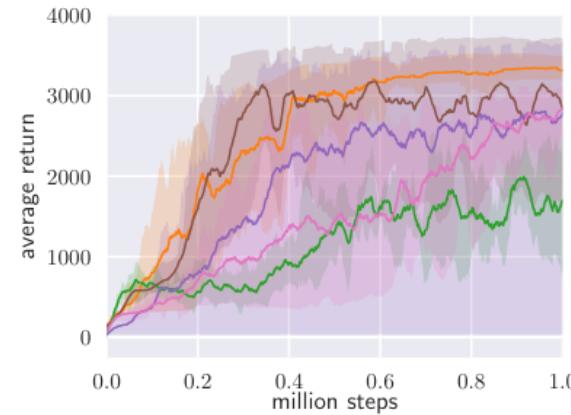
Advantages:

- Stable
- Learns many near-optimal policies
- Exploration
- Insensitivity to hyperparameters
- Off-Policy

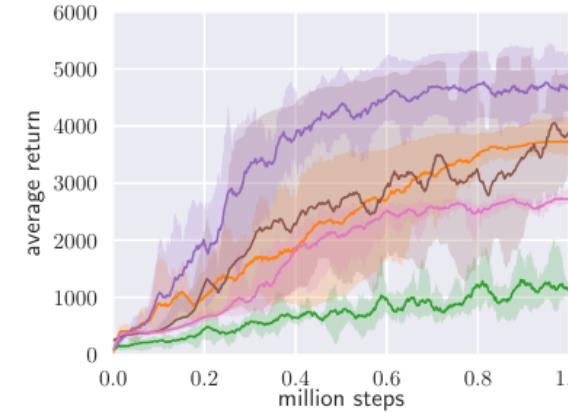
Soft Actor Critic

Advantages:

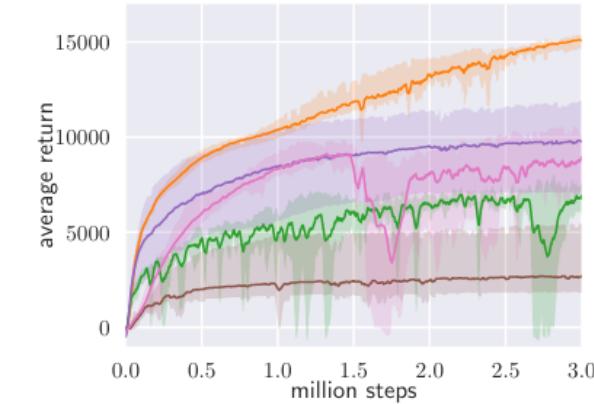
- Stable
- Learns |
- Explora
- Insensit
- Off-Poli



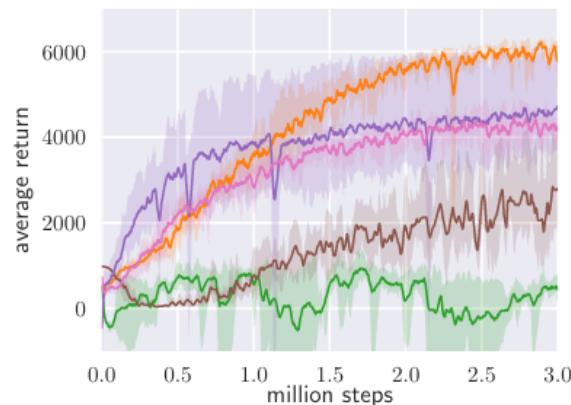
(a) Hopper-v1



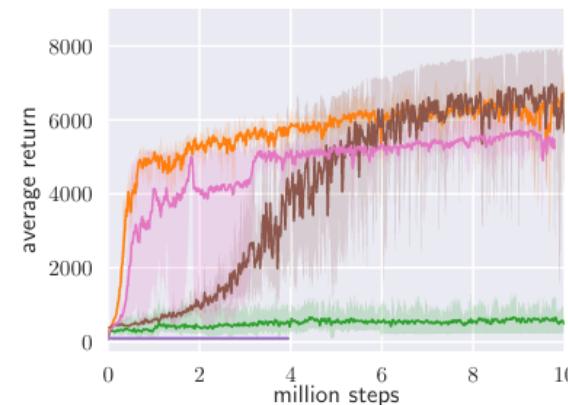
(b) Walker2d-v1



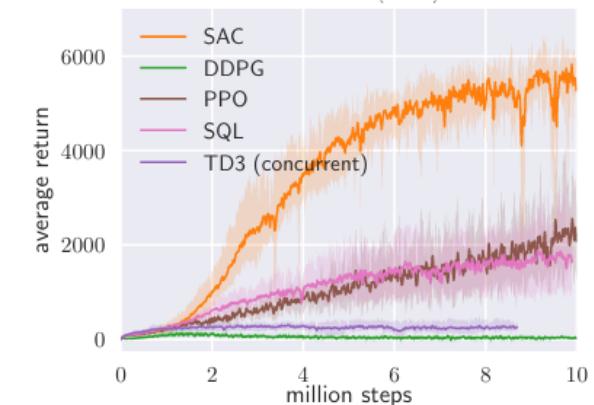
(c) HalfCheetah-v1



(d) Ant-v1



(e) Humanoid-v1



(f) Humanoid (rllab)

Soft Actor Critic

Advantages:

- Stable
- Learns many near-optimal policies
- Exploration
- Insensitivity to hyperparameters
- Off-Policy

Disadvantages

- Sensitive to α Solution = Entropy *constraint* and adjust α

Soft Actor Critic

Advantages:

- Starts with **Algorithm 1** Soft Actor-Critic

- Learns **Input:** θ_1, θ_2, ϕ

$$\theta_1 \leftarrow \theta_1, \theta_2 \leftarrow \theta_2$$

- Expects $\mathcal{D} \leftarrow \emptyset$

- Iterates **for each iteration do**

- for each environment step do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

for each gradient step do

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\alpha \leftarrow \alpha - \lambda \hat{\nabla}_\alpha J(\alpha)$$

$$\bar{\theta}_i \leftarrow \tau \theta_i + (1 - \tau) \bar{\theta}_i \text{ for } i \in \{1, 2\}$$

end for

end for

Output: θ_1, θ_2, ϕ

Disadvantages

• Complex

▷ Initial parameters

▷ Initialize target network weights

▷ Initialize an empty replay pool

▷ Sample action from the policy

▷ Sample transition from the environment

▷ Store the transition in the replay pool

▷ Update the Q-function parameters

▷ Update policy weights

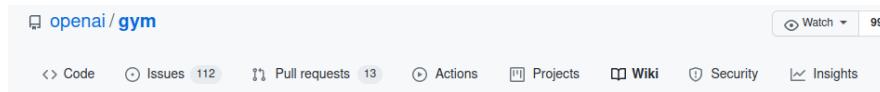
▷ Adjust temperature

▷ Update target network weights

▷ Optimized parameters

Wisdom

Deep RL: The Dream



Leaderboard

Aman Arora edited this page 3 days ago · 352 revisions

This page tracks the performance of user algorithms for various tasks in gym. Previously, users could submit their scores directly to gym.openai.com/envs, but it has been decided that a simpler wiki might do this task more efficiently.

This wiki page is a community driven page. Anyone can edit this page and add to it. We encourage you to contribute and modify this page and add your scores and links to your write-ups and code to reproduce your results. We also encourage you to add new tasks with the gym interface, but not in the core gym library (such as roboschool) to this page as well.

Links to videos are optional, but encouraged. Videos can be youtube, instagram, a tweet, or other public links. Write-ups should explain how to reproduce the result, and can be in the form of a simple gist link, blog post, or github repo.

We have begun to copy over the previous performance scores and write-up links over from the [previous page](#). This is an ongoing effort, and we can use some help.

Environments

Classic control

CartPole-v0

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The system is controlled by applying a force of +1 or -1 to the cart. The pendulum starts upright, and the goal is to prevent it from falling over. A reward of +1 is provided for every timestep that the pole remains upright. The episode ends when the pole is more than 15 degrees from vertical, or the cart moves more than 2.4 units from the center.



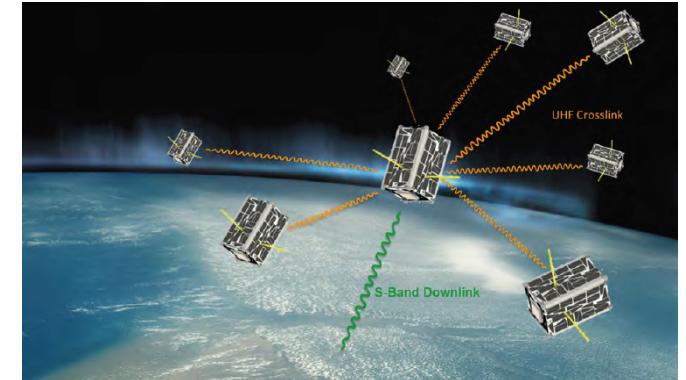
- [Environment Details](#)
- *CartPole-v0 defines "solving" as getting average reward of 195.0 over 100 consecutive trials.*
- *This environment corresponds to the version of the cart-pole problem described by Barto, Sutton, and Anderson [Barto83].*

| User | Episodes before solve | Write-up | Video |
|---------------|----------------------------------|------------------------------|-----------------------|
| Zhiqing Xiao | 0 (use close-form preset policy) | writeup | |
| Hengjian Jia | 0 (use close-form PID policy) | code/writeup | |
| Keavnn | 0 | writeup | |
| Shakti Kumar | 0 | writeup | Video |
| Nextgrid.ai 🌟 | 0 | writeup | Video |
| iRyanBell | 2 | writeup | |

Using Deep RL for your problem

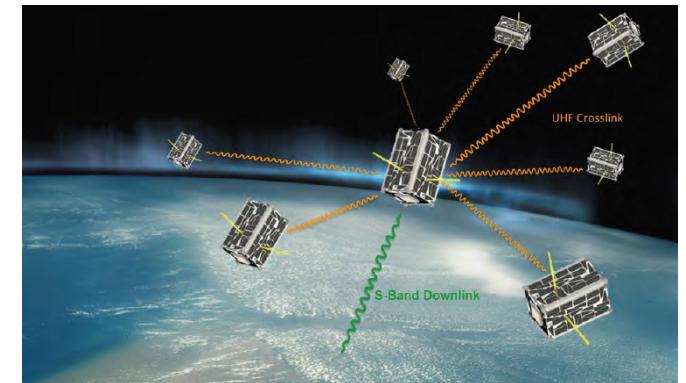
Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)



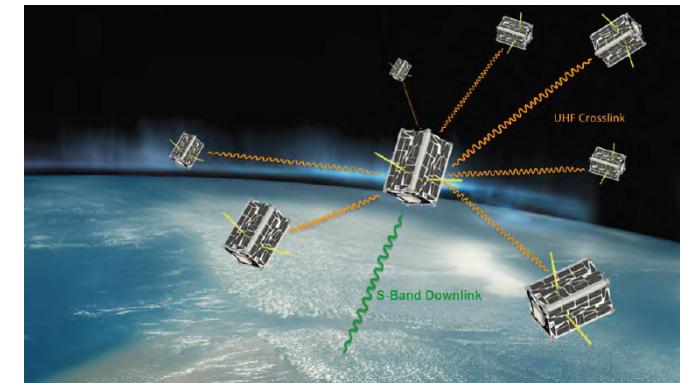
Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics



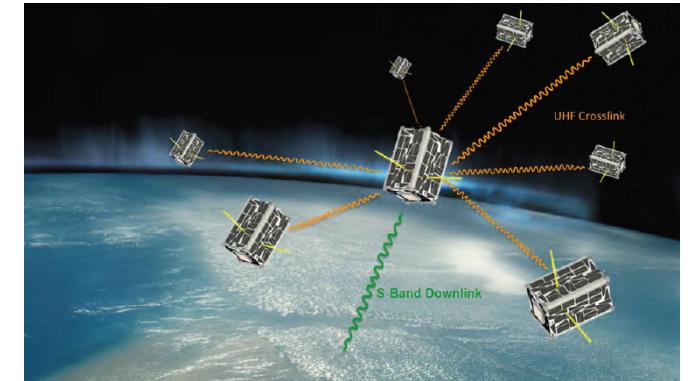
Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems



Using Deep RL for your problem

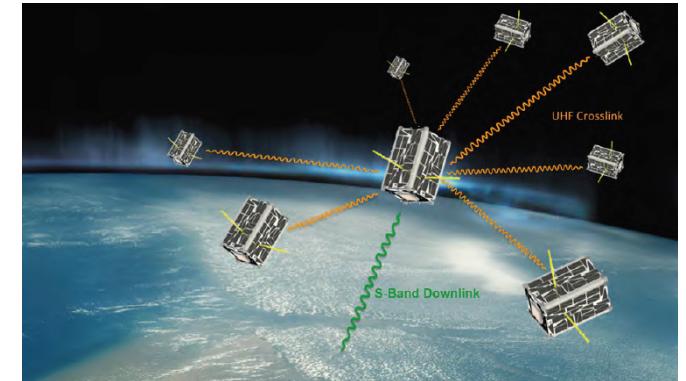
1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines



openai / **baselines**

Using Deep RL for your problem

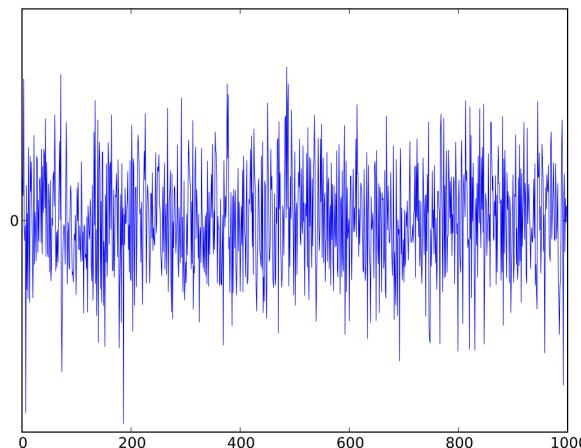
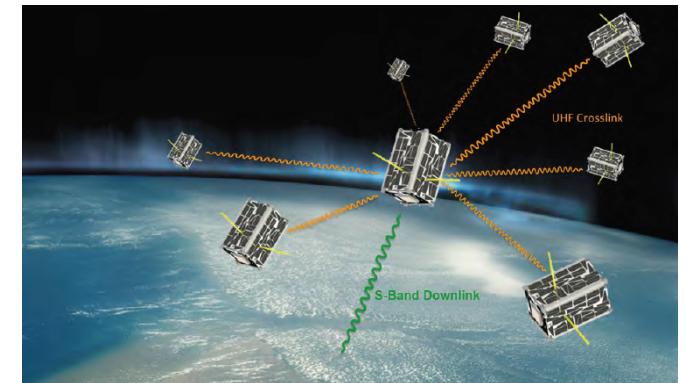
1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??



 [openai / baselines](#)

Using Deep RL for your problem

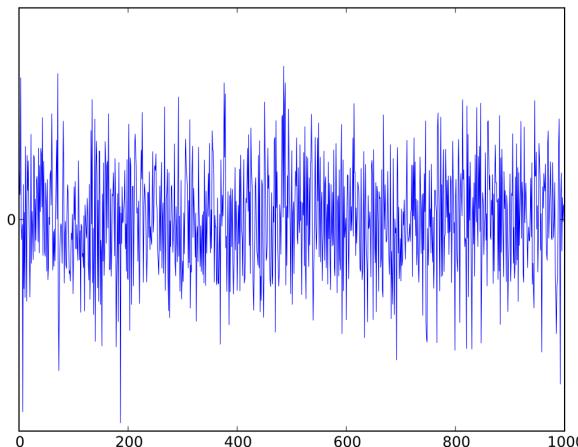
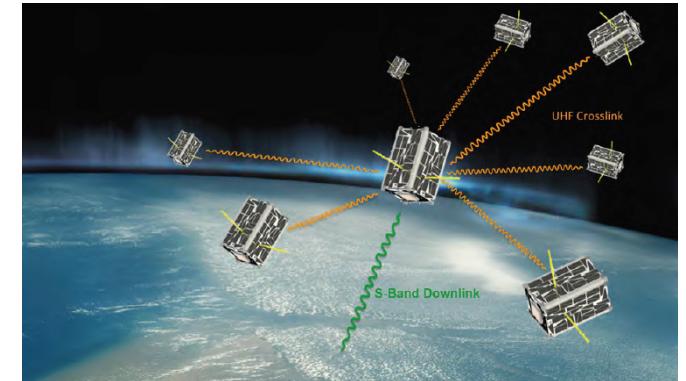
1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??



openai / **baselines**

Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??

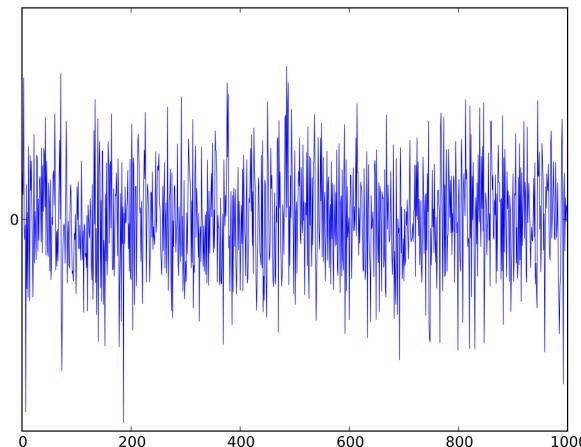
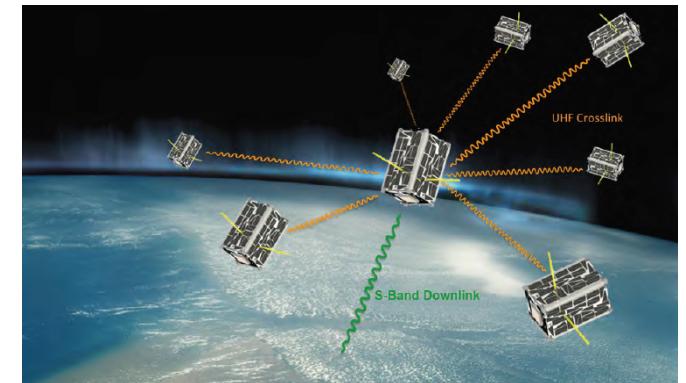


openai / baselines

Why not?

Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??



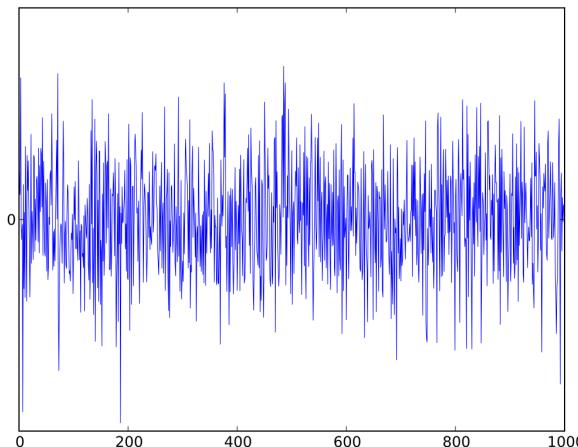
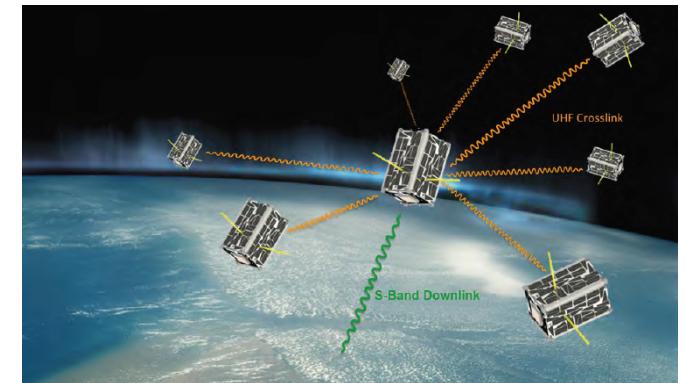
openai / baselines

Why not?

- Hyperparameters?

Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??



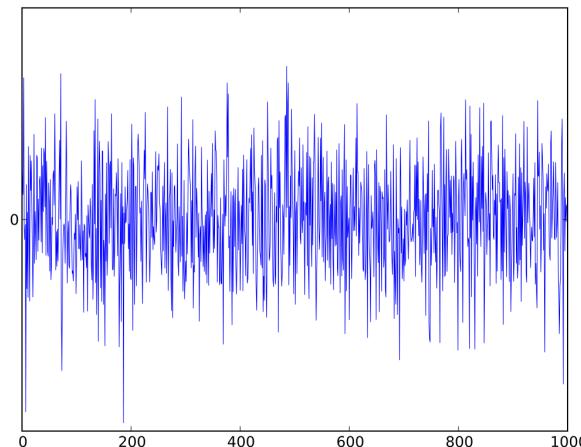
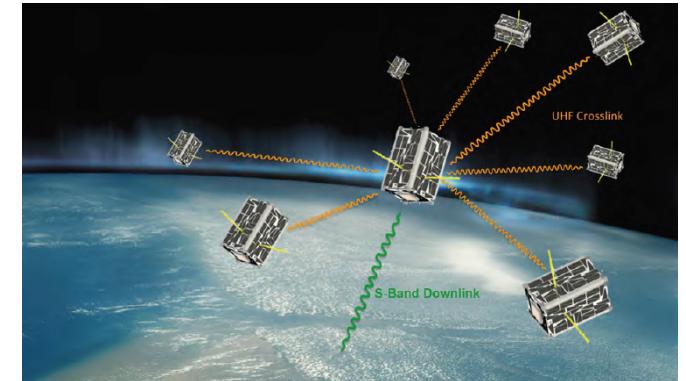
openai / baselines

Why not?

- Hyperparameters?
- Reward scaling?

Using Deep RL for your problem

1. Some interesting problem (smallsat swarm)
2. Spend weeks theorizing about the exact-right cost function and dynamics
3. Decide RL can solve all of your problems
4. Fire up open-ai baselines
5. Does it work??

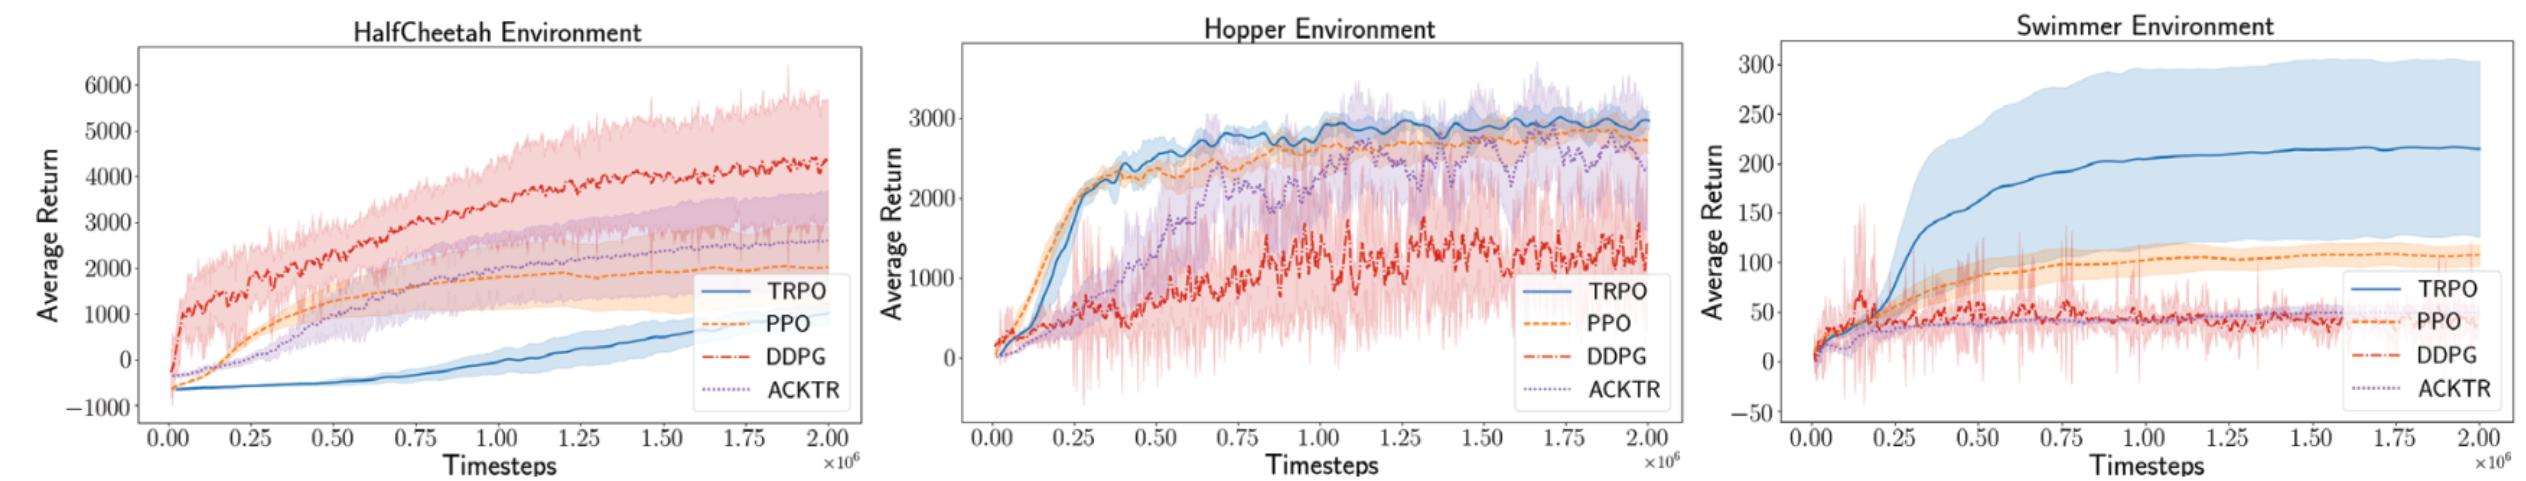


openai / baselines

Why not?

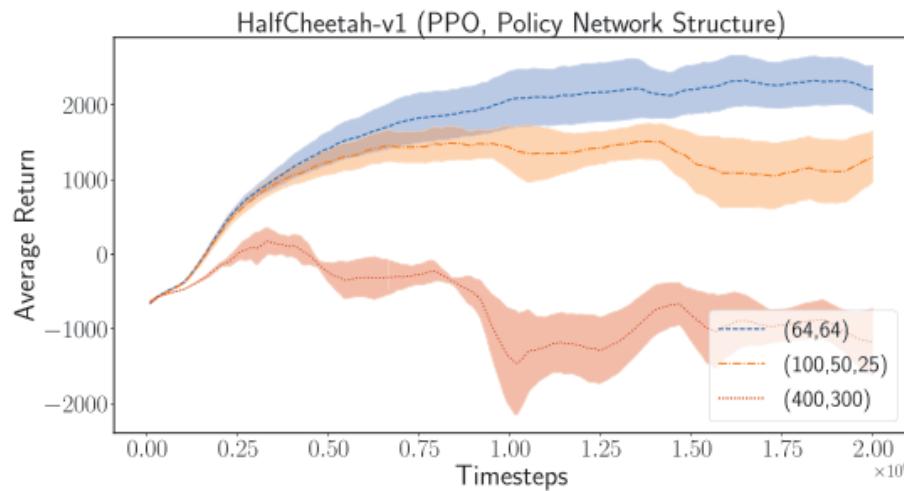
- Hyperparameters?
- Reward scaling?
- Not enough training time????

Algorithms

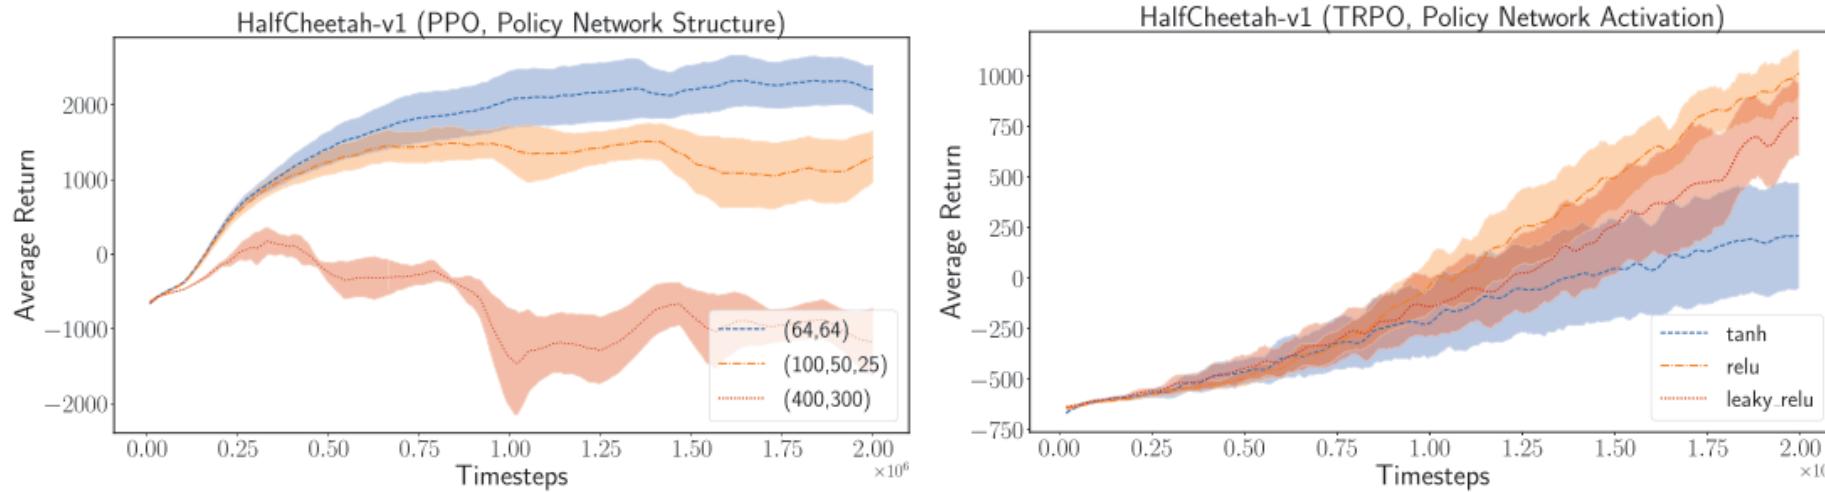


Policy Network Architecture

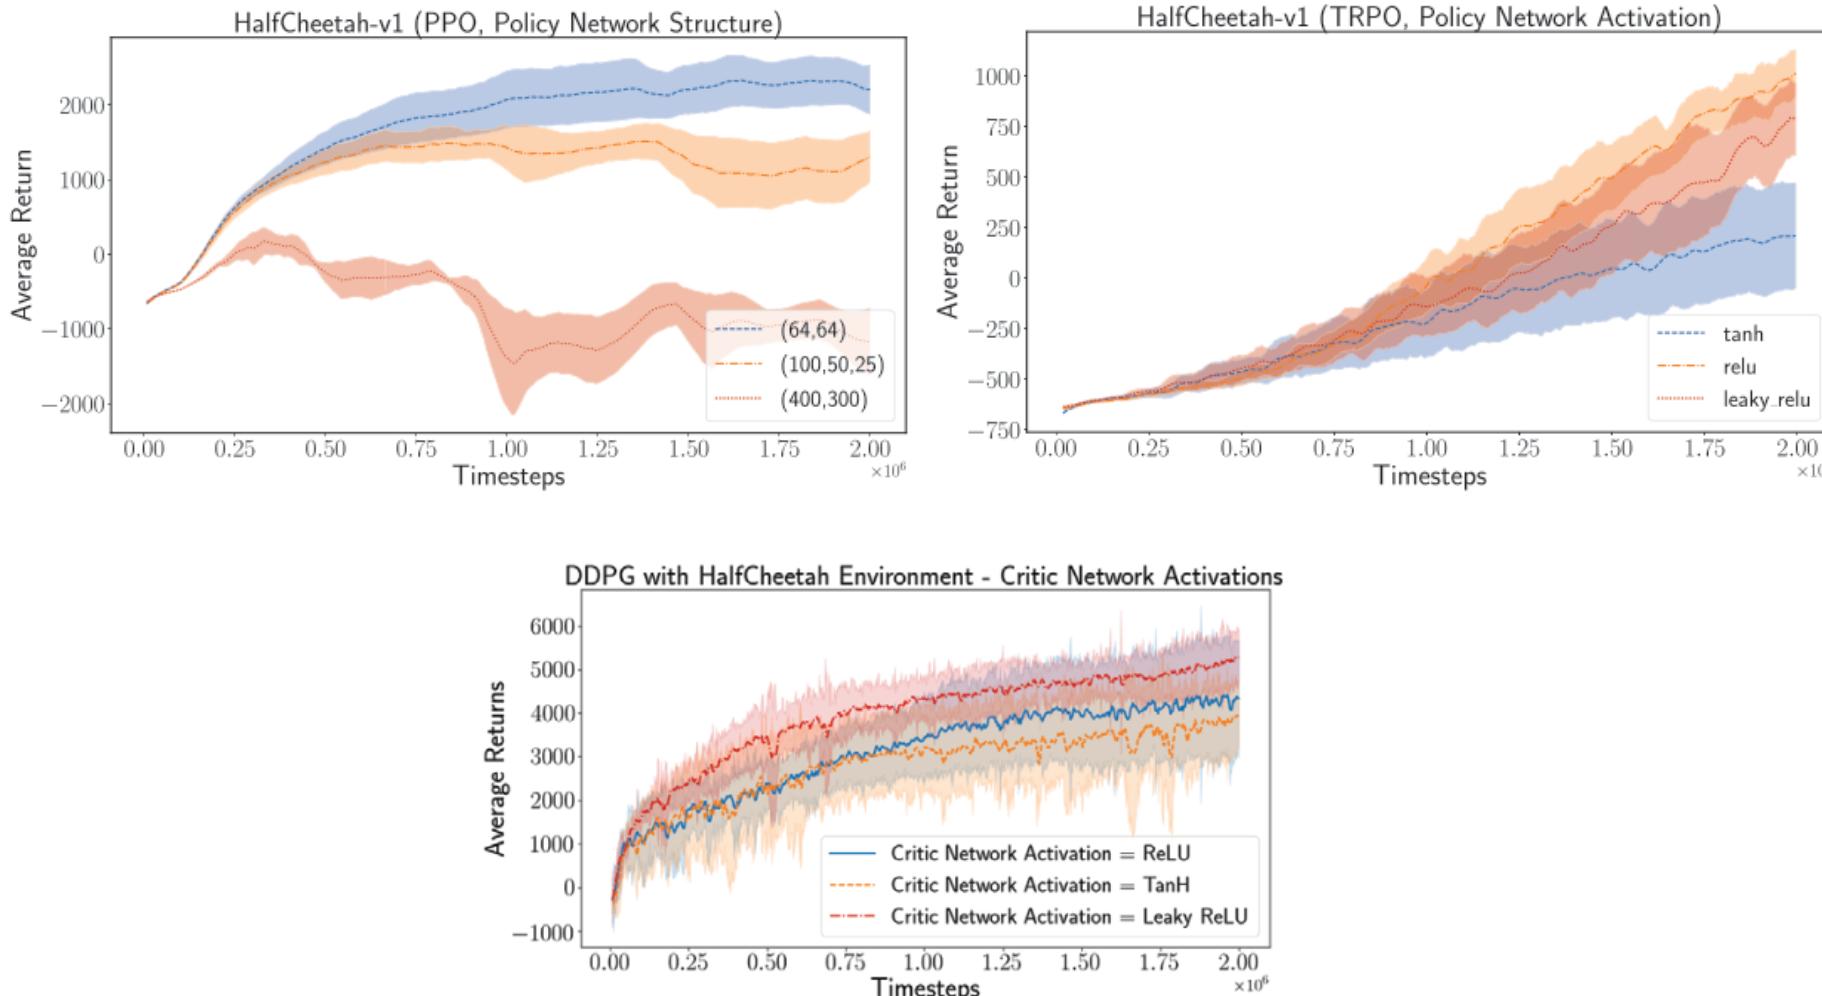
Policy Network Architecture



Policy Network Architecture



Policy Network Architecture



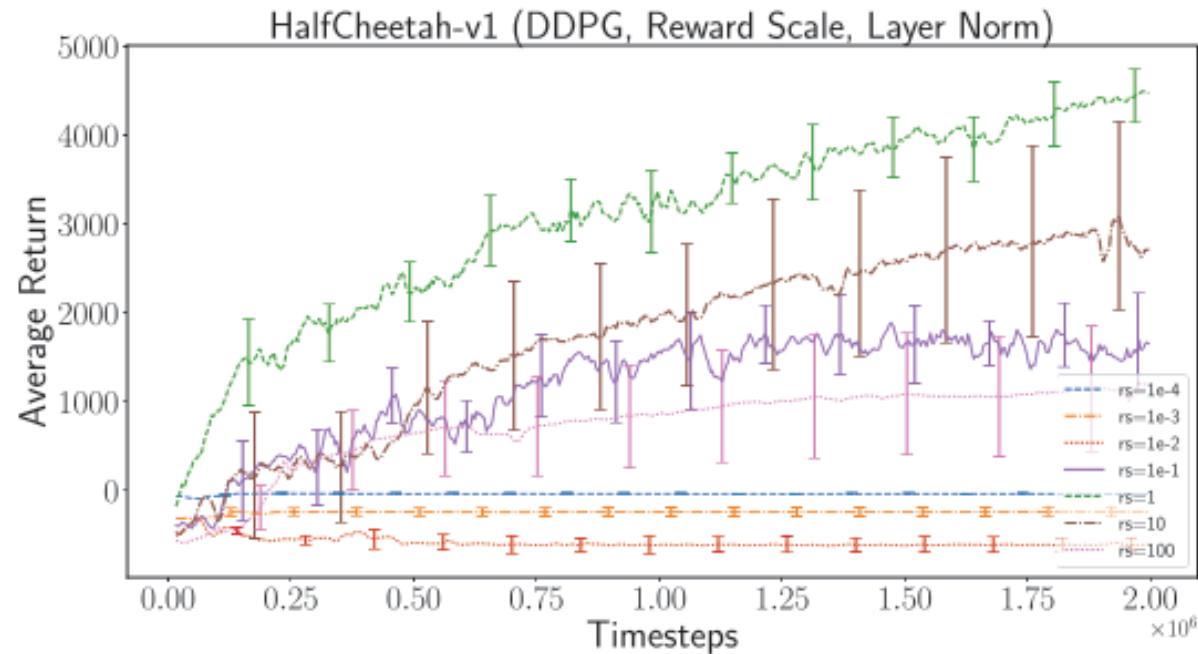
Reward Rescaling

Reward Rescaling

"simply multiplying the rewards generated from an environment by some scalar"

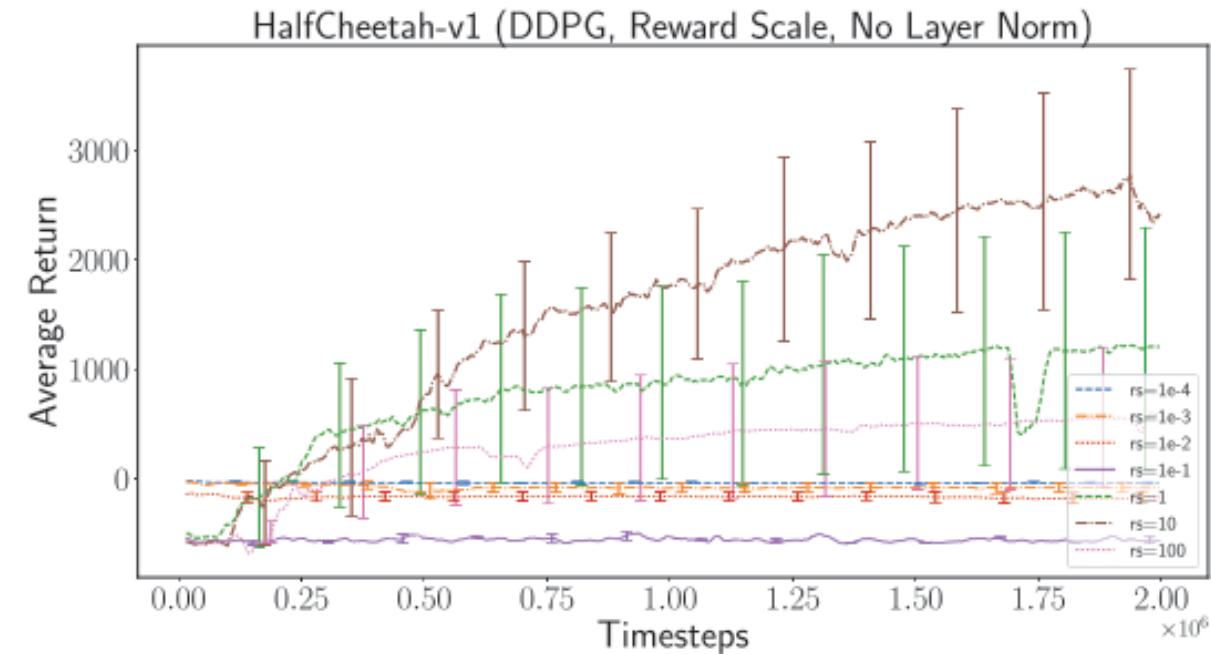
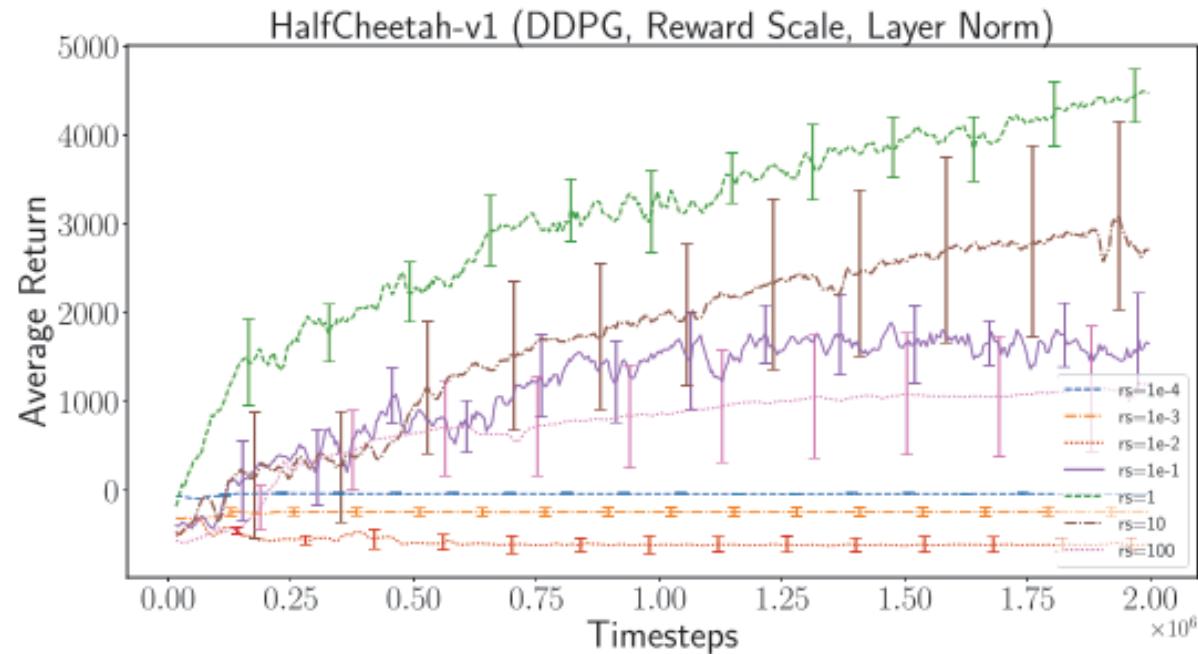
Reward Rescaling

"simply multiplying the rewards generated from an environment by some scalar"



Reward Rescaling

"simply multiplying the rewards generated from an environment by some scalar"



Statistical Significance

Statistical Significance

"Unfortunately, in recent reported results, it is not uncommon for the top-N trials to be selected from among several trials (Wu et al. 2017; Mnih et al. 2016)"

Statistical Significance

"Unfortunately, in recent reported results, it is not uncommon for the top-N trials to be selected from among several trials (Wu et al. 2017; Mnih et al. 2016)"

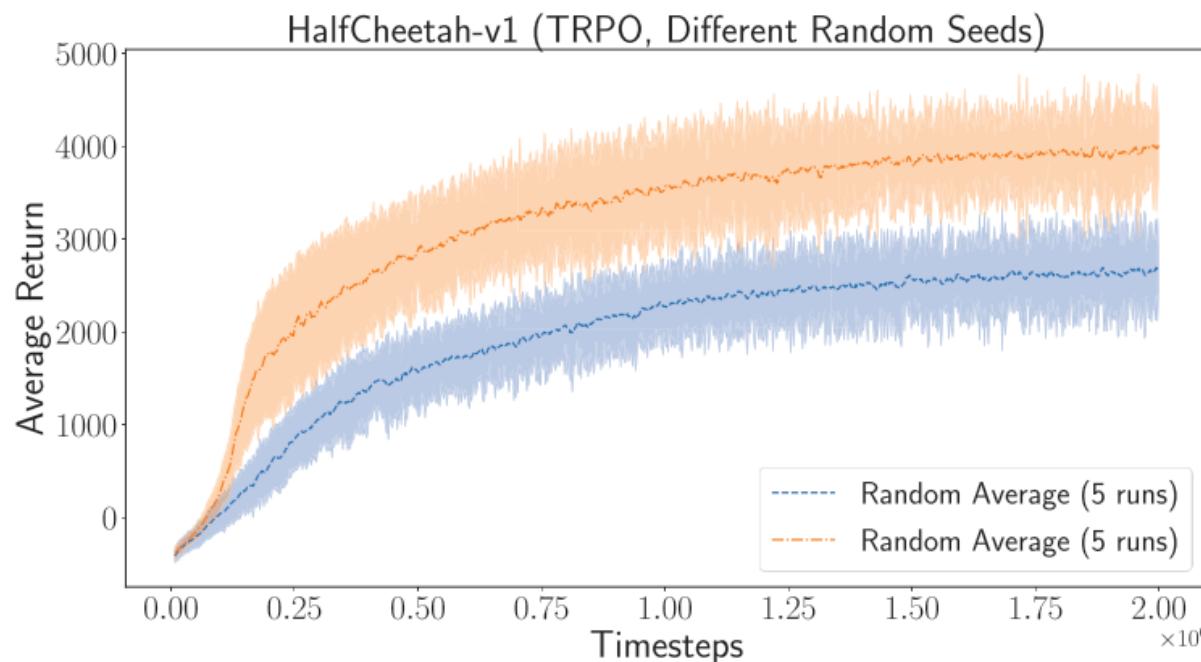
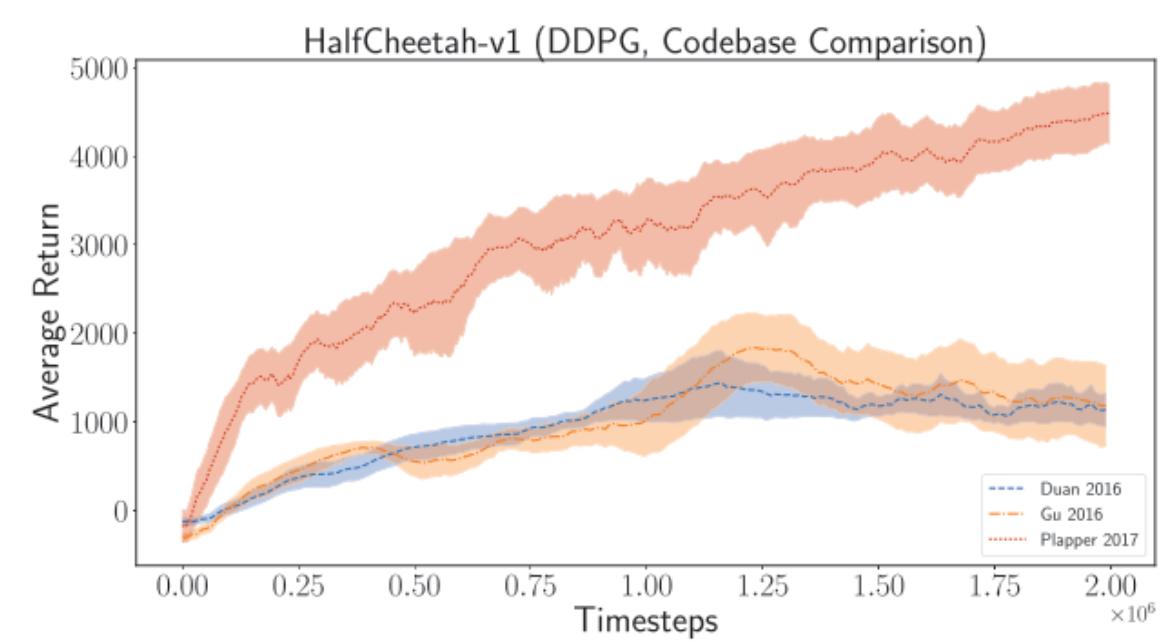
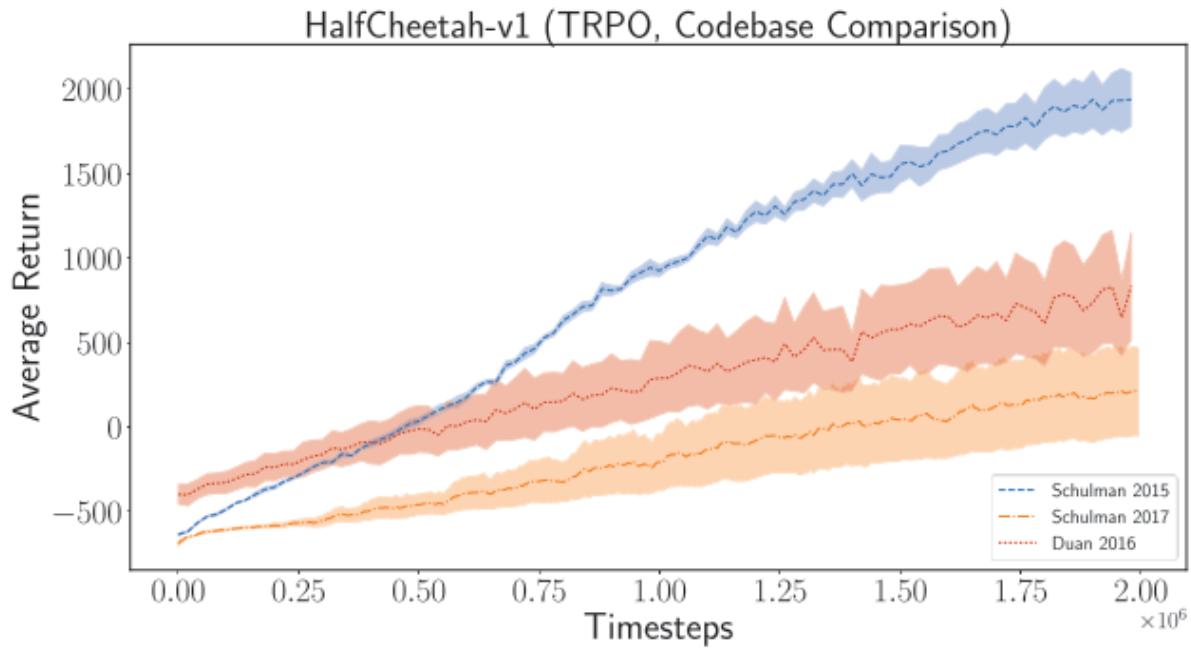


Figure 5: TRPO on HalfCheetah-v1 using the same hyperparameter configurations averaged over two sets of 5 different random seeds each. The average 2-sample t -test across entire training distribution resulted in $t = -9.0916$, $p = 0.0016$.

Codebases



How to choose an RL Algorithm

How to choose an RL Algorithm

(According to Sergey Levine)

How to choose an RL Algorithm

(According to Sergey Levine)



How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

Ease of Use
/ Stability

How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

← Fewer Samples

Ease of Use
/ Stability



How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

← Fewer Samples

Ease of Use
/ Stability

More Stable (?) →

How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

← Fewer Samples

Model-Based
RL

Ease of Use
/ Stability

More Stable (?) →

How to choose an RL Algorithm

(According to Sergey Levine)



Sample
Efficiency

← Fewer Samples

Model-Based
RL

Ease of Use
/ Stability

More Stable (?) →

Model-Based
Deep RL

How to choose an RL Algorithm

(According to Sergey Levine)



Sample Efficiency

← Fewer Samples

Model-Based RL

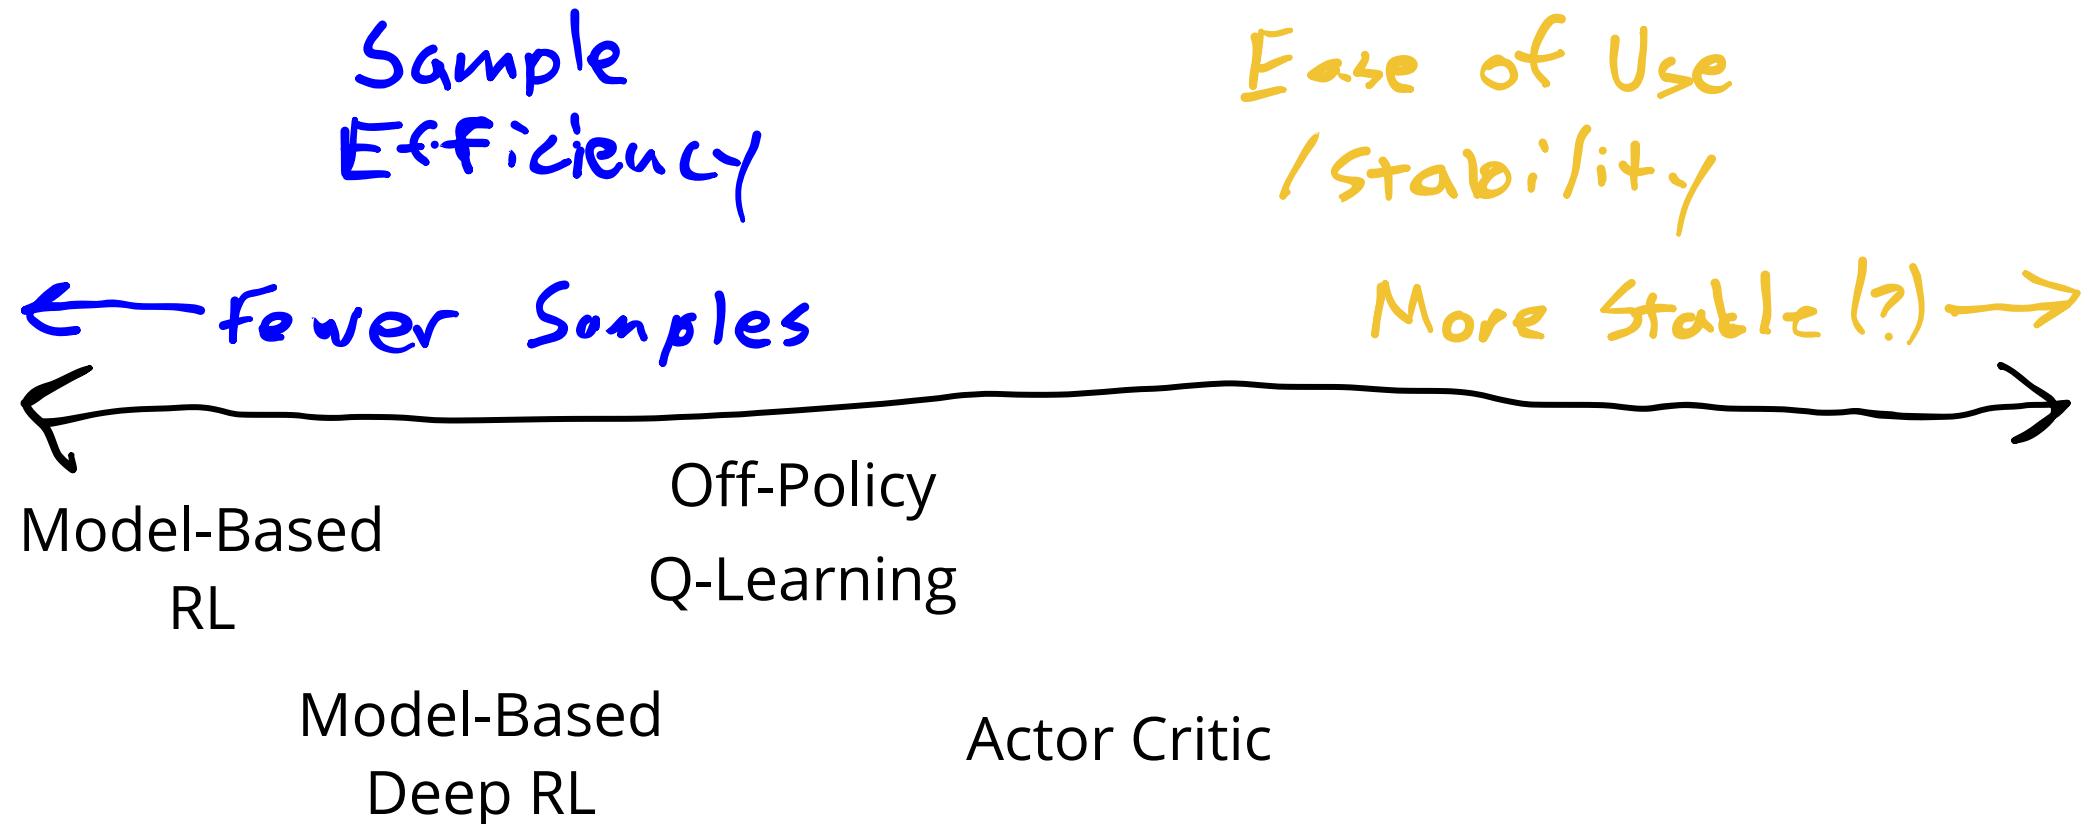
Model-Based Deep RL

Ease of Use / Stability

More Stable (?) →

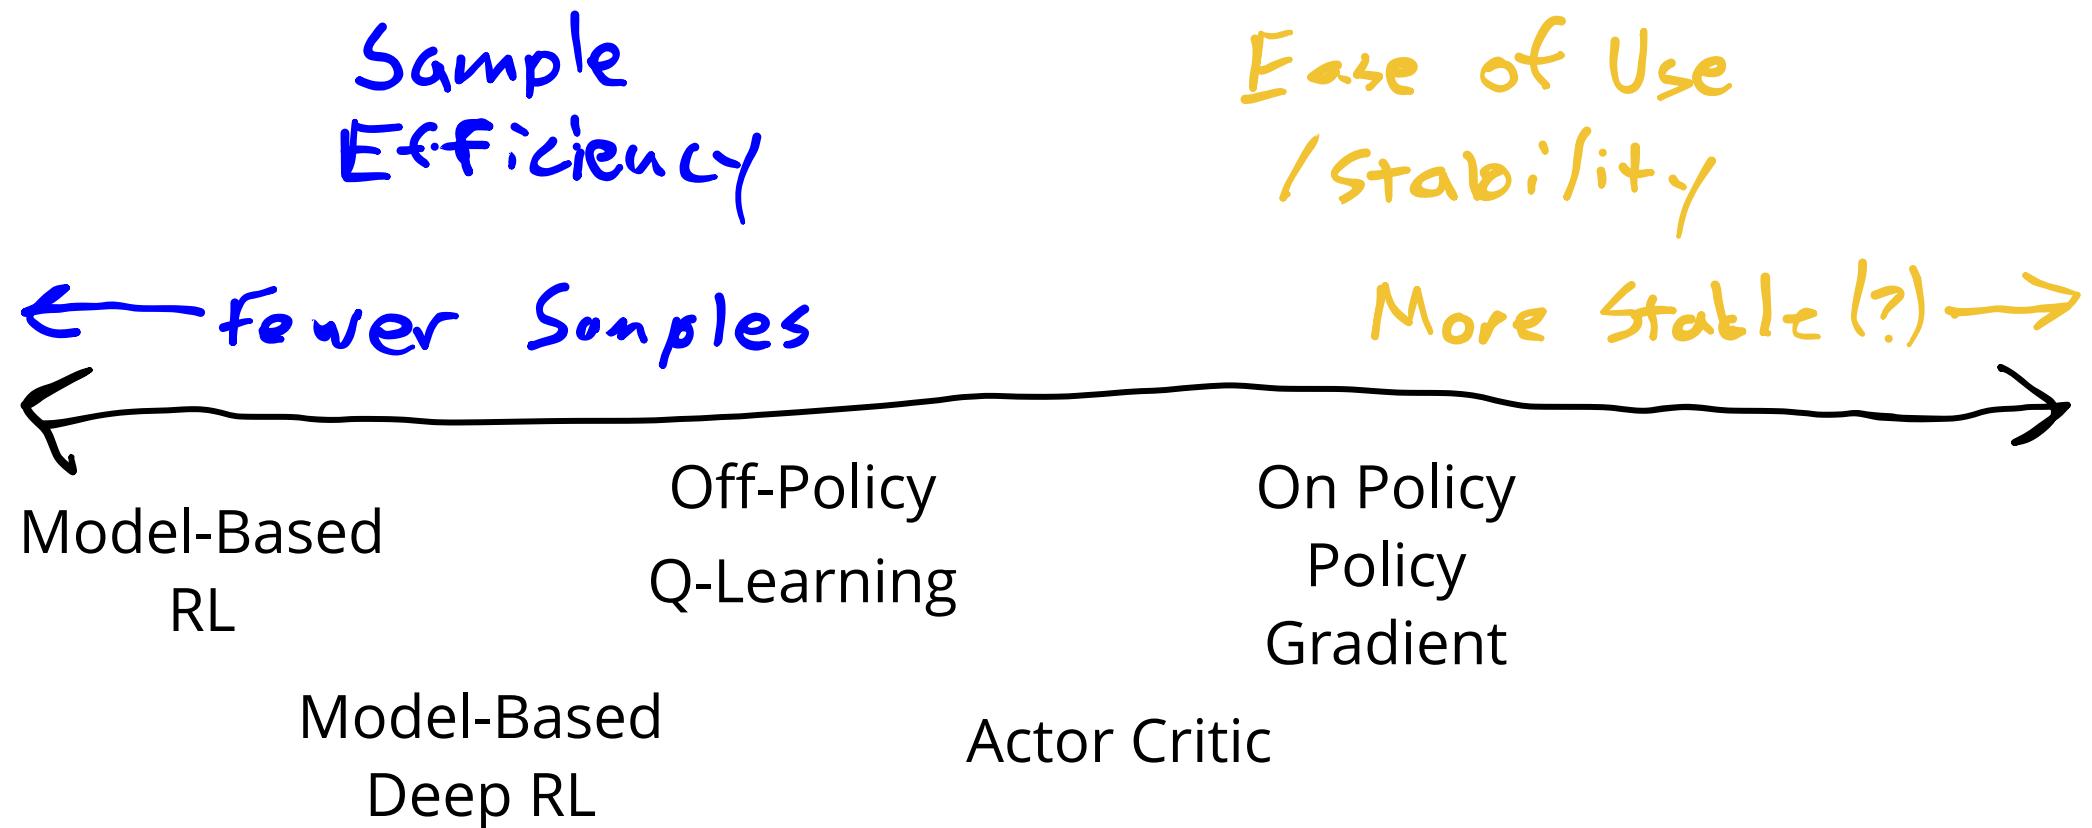
How to choose an RL Algorithm

(According to Sergey Levine)



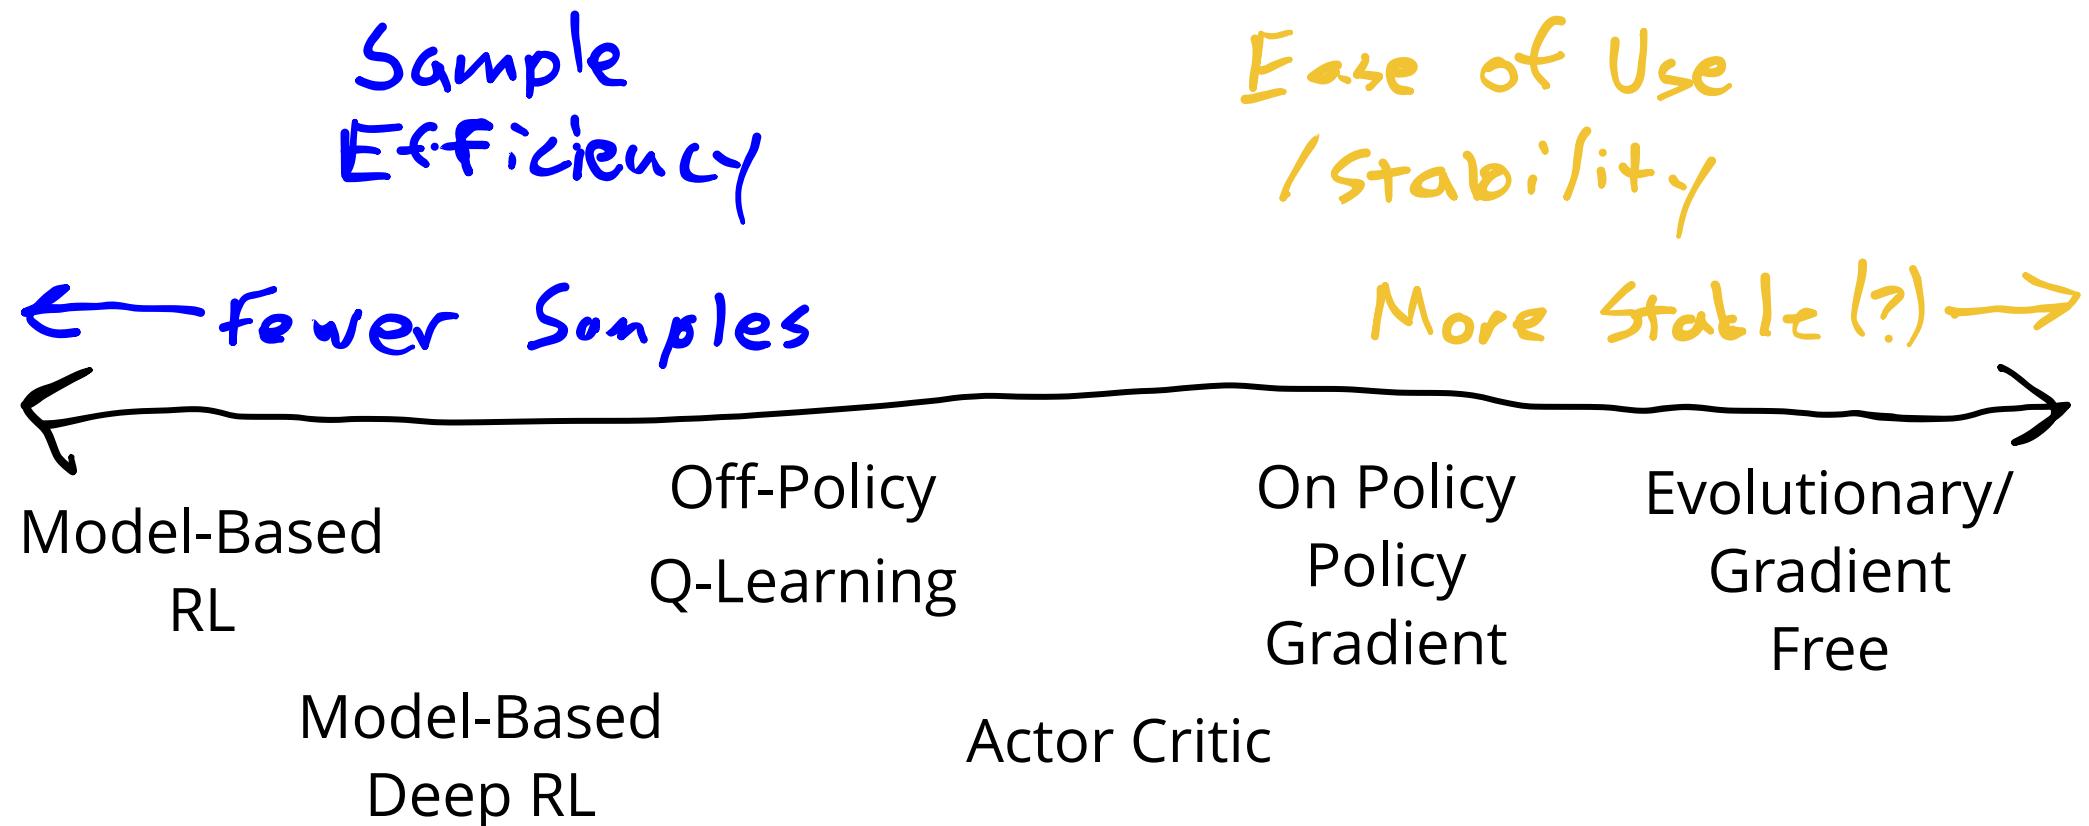
How to choose an RL Algorithm

(According to Sergey Levine)



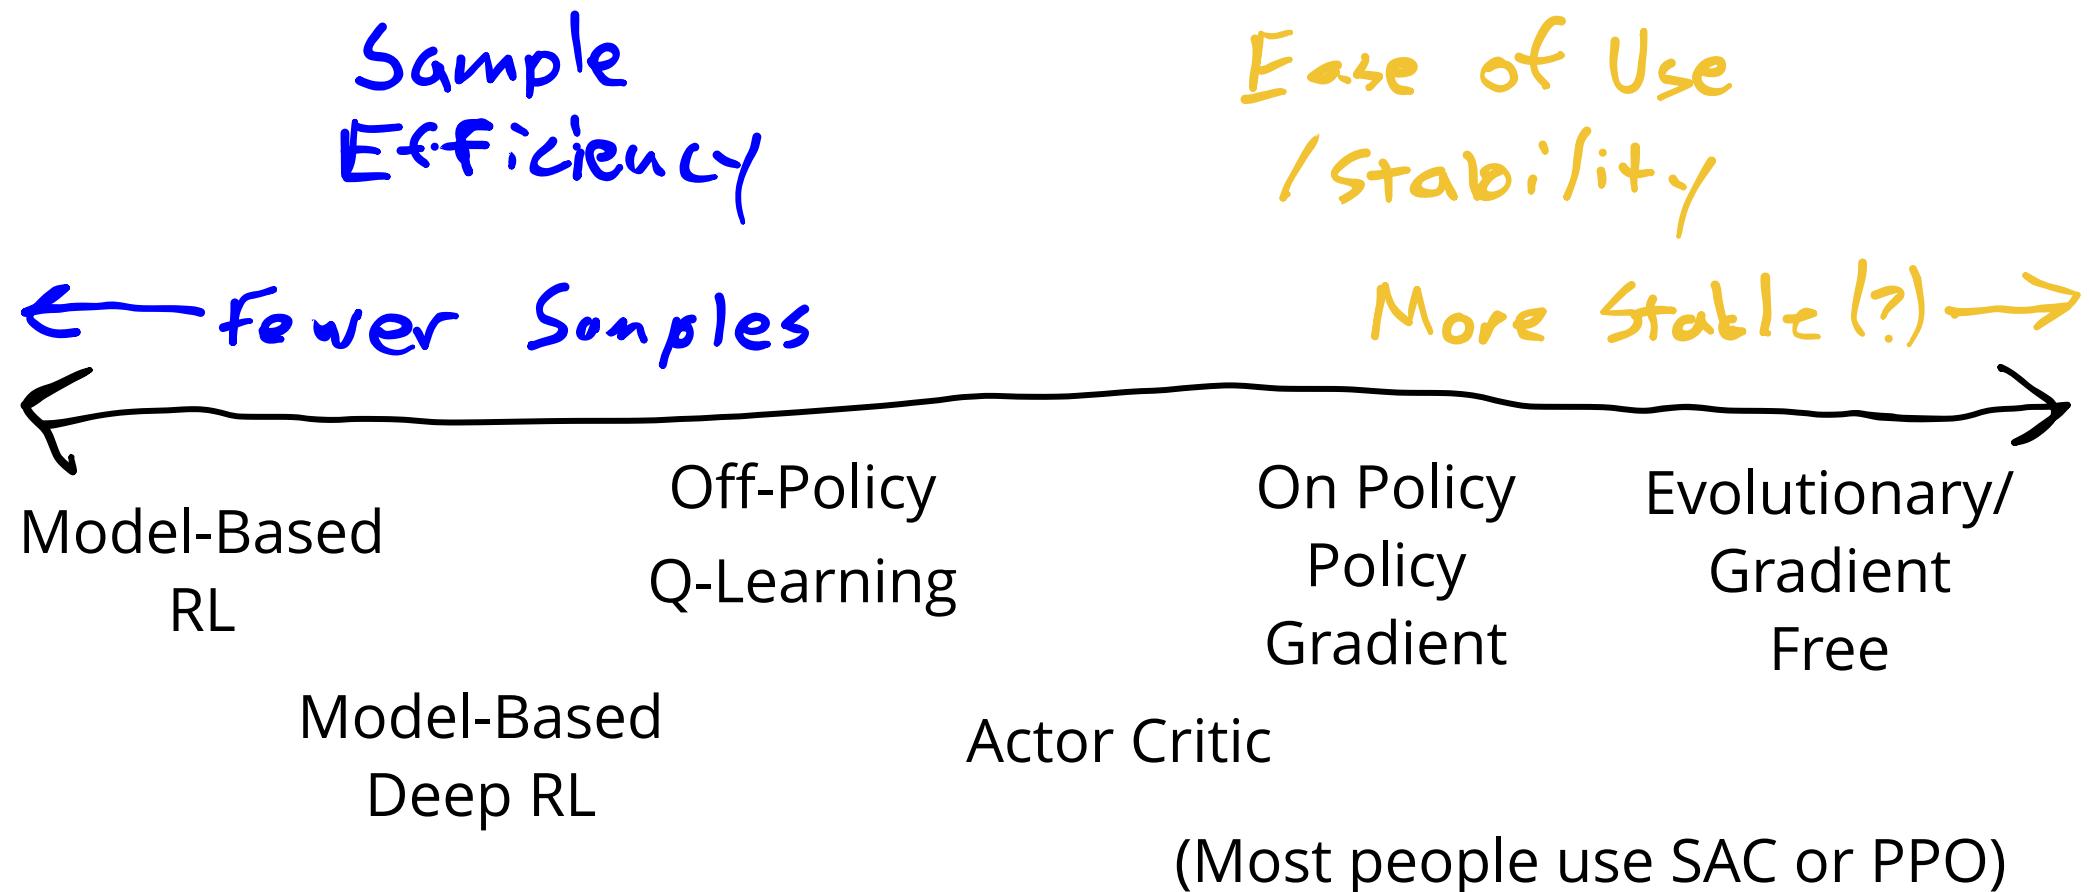
How to choose an RL Algorithm

(According to Sergey Levine)



How to choose an RL Algorithm

(According to Sergey Levine)



How to be successful with RL

- Always start with a small problem that works and scale up (keep verifying that it works with every change)
- Plot everything that you can think of (TensorBoard)
 - *Losses*
 - Policies
 - Value functions
 - Trajectories
 - (Average return) Learning curve
- Keep calm and lower your learning rate

